



CS590 - Algorithms

Lecture 6 – Binary Search Trees

Close your laptop



6. Binary Search Trees (BST)

6.1. Binary Search Trees

6.2. BST – In order tree walk

6.3. Querying a BST

6.4. Insertion

6.5. Deletion


6.6. Expected Height of a Randomly Built BST



6.1. Binary search trees

- **Binary search trees (BSTs)** are an important data structure for dynamic sets.
- They accomplish many dynamic set operations in $O(h)$ time, where h = height of tree.
- We represent a **binary tree** by a linked data structure in which each node is an object.
- $T.root$ points to the root of tree T .
- Each node contains the fields
 - *Key*: (and possibly other satellite data).
 - *left*: points to left child.
 - *right*: points to right child.
 - *p*: points to parent. $\Rightarrow T.root.p = \text{NIL}$.
- Stored keys must satisfy the **binary-search-tree property**.
 - If y is in left subtree of x , then $y.key \leq x.key$.
 - If y is in right subtree of x , then $y.key \geq x.key$.

6.1. Binary search trees



The *binary-search-tree property* allows us to print keys in a binary search tree in order, recursively, using an algorithm called an *inorder-tree-walk*. Elements are printed in monotonically increasing order.

6.2. BST – In order tree walk


Algorithm (INORDER-TREE-WALK(x))

```
(1)  if ( $x \neq NIL$ ) then
(2)    INORDER-TREE-WALK( $x$ .left)
(3)    print  $x$ .key
(4)    INORDER-TREE-WALK( $x$ .right)
(5)  fi
```

How INORDER-TREE-WALK works?

- Check to make sure that x is not NIL.
- Recursively, print the keys of the nodes in x 's left subtree.
- Print x 's key.
- Recursively, print the keys of the nodes in x 's right subtree.

6.2. BST – In order tree walk



BST prints: A, B, D, F, H, K

Correctness: Follows by induction directly from the binary-search-tree property.

Time: Intuitively, the walk takes $\Theta(n)$ time for a tree with n nodes, because we visit and print each node once.



6.2. BST – In order tree walk

Recursion Running Time Equation Construction:

- Let $T(n)$ denote the time taken by INORDER-TREE-WALK when it is called on the root of an n -node subtree.
- We have $T(n) = \Omega(n)$ since it visits all n nodes of the subtree.
- The amount of time on an empty subtree is $T(0) = c$ for $c > 0$.
- **Inductive Step:**
 - For $n > 0$, suppose that a node x on the left subtree has k nodes.
 - Then the right subtree has $n - k - 1$ nodes.
 - The performance time is bounded by $T(n) \leq T(k) + T(n - k - 1) + d$ for some constant $d > 0$.
- We need to show that $T(n) = O(n)$!



6.2. BST – In order tree walk

- $T(n) \leq T(k) + T(n - k - 1) + d$ for some constant $d > 0$.
- Using the substitution method with guessing that $T(n) \leq (c + d)n + c$.
- For $n = 0$, we see that $T(0) = c$.
- For $n > 0$, we have

$$\begin{aligned} T(n) &\leq T(k) + T(n - k - 1) + d \\ &\leq ((c + d)k + c) + ((c + d)(n - k - 1) + c) + d \\ &= (c + d)n + c - (c - d) + c + d \\ &= (c + d)n + c \end{aligned}$$



6.3. Querying a BST

Searching

- Need to search for a key sorted in BST.
- BST can support queries such as MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR.

Algorithm (TREE-SEARCH(x, k))

```
(1)  if ( $x = NIL \mid k = x.key$ ) then
(2)    return  $x$ 
(3)  if ( $k < x.key$ ) then
(4)    return TREE-SEARCH( $x.left, k$ )
(5)  else
(6)    return TREE-SEARCH( $x.right, k$ )
```

- Initial call is TREE-SEARCH($T.root, k$).
- **Time:** The algorithm recurses, visiting nodes on a downward path from the root. Thus, running time is $O(h)$, where h is the height of the tree.



6.3. Querying a BST

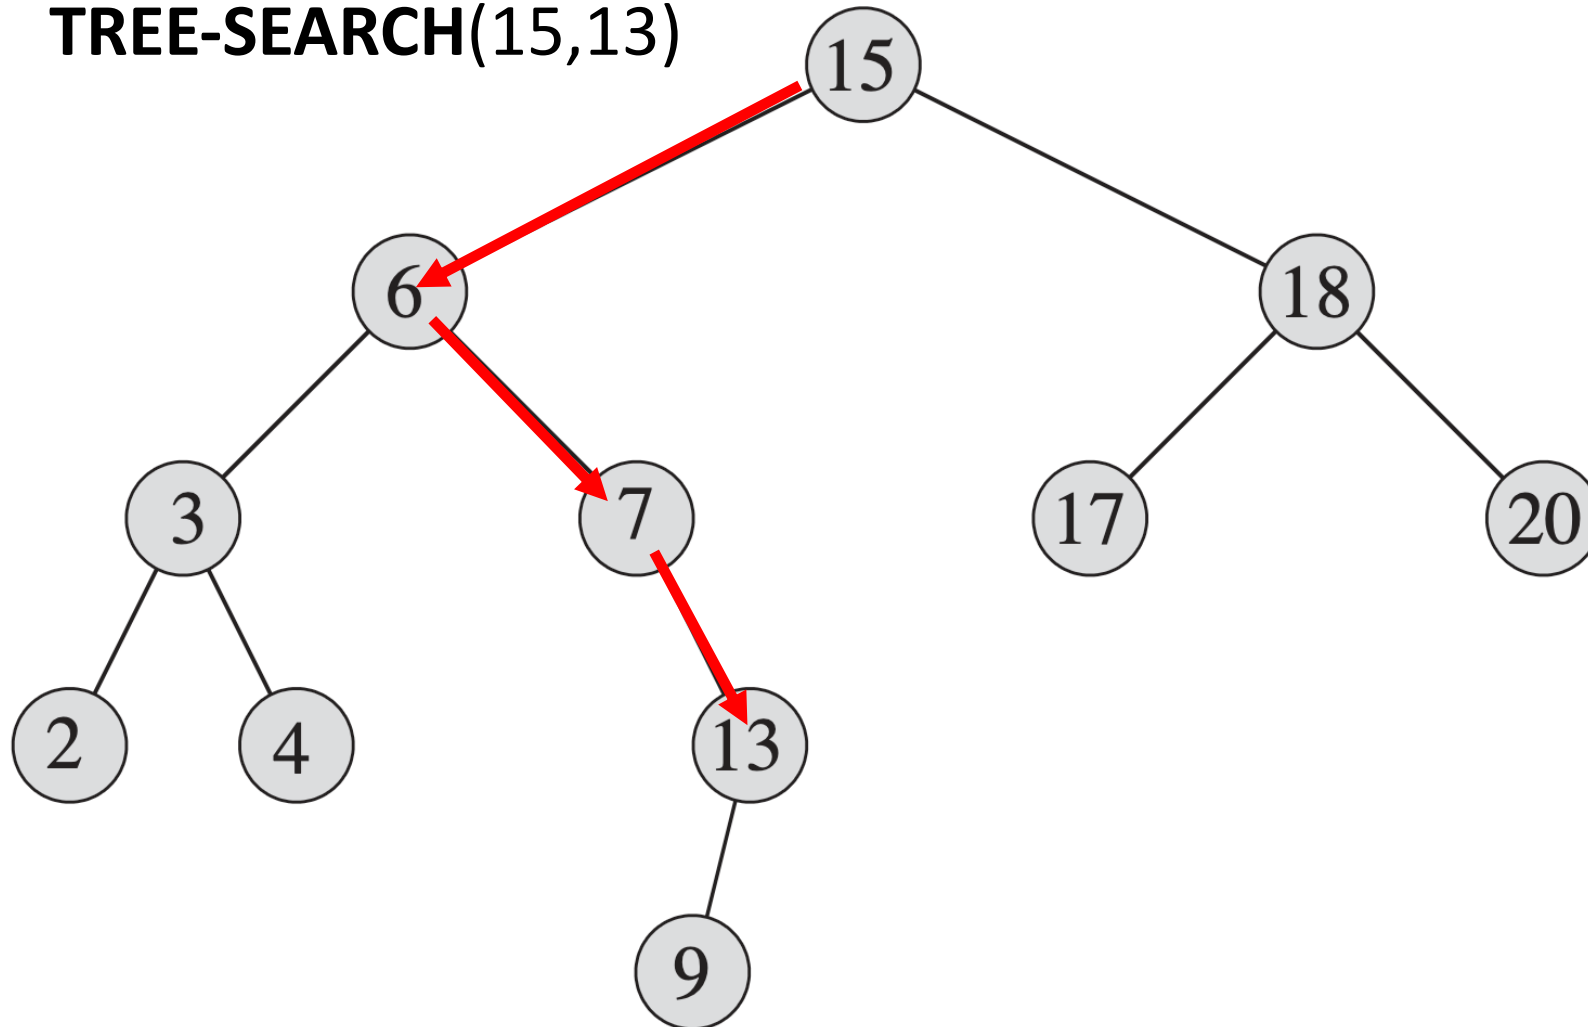
Algorithm (ITERATIVE-TREE-SEARCH(x, k))

```
(1)  while  $x \neq NIL$  &  $k \neq x.key$ 
(2)    if  $k < x.key$ 
(3)       $x = x.left$ 
(4)    else  $x = x.right$ 
(5)  return  $x$ 
```

6.3. BST – Tree Search



TREE-SEARCH(15,13)



6.3. Querying a BST

Minimum and maximum

Algorithm (TREE-MINIMUM(x))

```
(1)  while  $x.\text{left} \neq \text{NIL}$  do
(2)       $x = x.\text{left}$ 
(3)  return  $x$ 
```

Algorithm (TREE-MAXIMUM(x))

```
(1)  while  $x.\text{right} \neq \text{NIL}$  do
(2)       $x = x.\text{right}$ 
(3)  return  $x$ 
```

The BST property guarantees that


- The minimum key of a BST is located at the leftmost node.
- The maximum key of a BST is located at the rightmost node.
- **Time:** Both procedures visit nodes that form a downward path from the root to a leaf.
- Both procedures run in $O(h)$ time, where h is the height of the tree.
- Traverse the appropriate pointers (*left* or *right*) until NIL is reached.

6.3. Querying a BST

Successor and predecessor

Assuming that all keys are distinct,

- the **successor** of a node x is the node y such that $y.key$ is the smallest $key > x.key$.
- the **predecessor** of a node x is the node y such that $y.key$ is the largest $key < x.key$.





6.3. Querying a BST

Successor and predecessor

- We can find x 's successor based entirely on the tree structure.
- No key comparisons are necessary.
- If x has the largest /smallest key in the binary search tree, then successor/predecessor is NIL.
- There are two cases:
 1. If node x has a non-empty right subtree, then x 's successor is the minimum in x 's right subtree.
 2. If node x has an empty right subtree, notice that:
 - As long as we move to the left up the tree (move up through right children), we're visiting smaller keys.
 - x 's successor y is the node that x is the predecessor of y (x is the maximum in y 's left subtree).

6.3. Querying a BST

Successor and predecessor

Algorithm (TREE-SUCCESSOR(*x*))

```
(1)  if x.right  $\neq$  NIL then
(2)    return TREE-MINIMUM(x.right)
(3)  y = x.p
(4)  while (y  $\neq$  NIL and x = y.right) do
(5)    x = y
(6)    y = y.p
(7)  return y
```

6.3. Querying a BST

Successor and predecessor


Algorithm (TREE-PREDECESSOR(x))

```
(1)  if  $x.\text{left} \neq \text{NIL}$  then
(2)    return TREE-MAXIMUM( $x.\text{left}$ )
(3)   $y = x.p$ 
(4)  while ( $y \neq \text{NIL}$  and  $x = y.\text{left}$ ) do
(5)     $x = y$ 
(6)     $y = y.p$ 
(7)  return  $y$ 
```

- **Time:** For both the TREE-SUCCESSOR and TREE-PREDECESSOR procedures, in both cases, we visit nodes on a path down the tree or up the tree.
- Thus, running time is $O(h)$, where h is the height of the tree.

6.3. Querying a BST

Successor and predecessor



- Find the successor of the node with key value 15.
- Find the successor of the node with key value 6.
- Find the successor of the node with key value 4.
- Find the predecessor of the node with key value 6.



6.4. Insertion

- Insertion and deletion allows the dynamic set represented by a binary search tree to change.
- The binary-search-tree property must hold after the change.
- To insert value v into the binary search tree, we create a node z , with $z.key=v$, $z.left=NIL$, and $z.right=NIL$.
- Find the position for z by tracing downward path from the root. Two points must be maintained:
 - Pointer x : traces downward path.
 - Pointer y : “trailing pointer” to keep track of parent of x .
- Traverse downward by comparing $x.key$ with v (or $z.key$) then move to the left or right child accordingly.
- The correct position for z if $x=NIL$.
- Compare z 's value with y 's value, and insert z at either y 's *left* or *right*, appropriately.

6.4. Insertion and Deletion




Algorithm (TREE-INSERT(T, x))

```
(1)  y = NIL, x = T.root
(2)  while (x  $\neq$  NIL) do
(3)    y = x
(4)    if (z.key < x.key) then
(5)      x = x.left
(6)    else x = x.right
(7)  z.p = y
(8)  if (y = NIL) then
(9)    T.root = z
(10) else if (z.key < y.key) then
(11)   y.left = z
(12) else y.right = z
```

6.4. Insertion and Deletion




TREE-INSERT (T, C)



6.4. Insertion and Deletion



TREE-INSERT (T, C)



Start: $y = \text{NIL}$, $x = F$

1. $x = F$, $y = F$
2. $C < F$: $x = B$
3. $y = B$, $C < B$, $x = D$
4. $y = D$, $C < D$, $x = \text{NIL}$
5. While-loop terminates: $y = D$, $x = \text{NIL}$


Algorithm (TREE-INSERT(T,x))

- (1) $y = \text{NIL}$, $x = T.\text{root}$
- (2) while ($x \neq \text{NIL}$) do
- (3) $y = x$
- (4) if ($x.\text{key} < C.\text{key}$) then
- (5) $x = x.\text{left}$
- (6) else $x = x.\text{right}$

6.4. Insertion and Deletion



TREE-INSERT (T, C)



$z.p = D$

$y = D, C < D, D.left = C$

Algorithm (TREE-INSERT(T,x))

(1) ... (6)

(7) $z.p = y$

(8) if ($y = N/L$) then

(9) $T.root = z$

(10) else if ($z.key < y.key$) then


(11) $y.left = z$

(12) else $y.right = z$

6.4. Insertion and Deletion



TREE-INSERT (T, C)



6.4. Insertion

- TREE-INSERT can be used with INORDER-TREE-WALK to sort a given set of numbers.

Algorithm (TREE-SORT(A))

```
(1)  let  $T$  be an empty binary search tree
(2)  for  $i \leftarrow 1$  to  $n$ 
(3)    do TREE-INSERT( $T, A[i]$ )
(4)  INORDER-TREE-WALK( $root[T]$ )
```

- Worst case: $\Theta(n^2)$ occurs when a linear chain of nodes results from the repeated TREE-INSERT operations.
- Best case: $\Theta(n \lg n)$ occurs when a binary tree of height $\Theta(\lg n)$ results from the repeated TREE-INSERT operations.



6.5. Deletion

For deletion, consider three cases:

Case 1: z has no children.

- Delete z by making the parent of z point to NIL, instead of to z .

Case 2: z has one child.

- Delete z by making the parent of z point to z 's child, instead of to z .

Case 3: z has two children.

- Find the successor of z , y , and replace z by the successor.
 - y must be in the right subtree of z and have no left child.
 - The rest of the original subtree of z becomes the new right subtree of y .
 - The left subtree of z becomes the new left subtree of y .

6.5. Deletion


An algorithm **TRANSPLANT** replaces one subtree as the child of its parent by another subtree.

```
Algorithm (TRANSPLANT (T, u, v ))
(1)   if (u.p = NIL) then
(2)     T.root = v
(3)   else
(4)     if (u = u.p.left) then
(5)       u.p.left = v
(6)     else u.p.right = v
(7)     If (v ≠ NIL) then
(8)       v.p = u.p
```


- Replaces the subtree rooted at u by the subtree rooted at v.
- Makes u.p becomes v
 - If u is the root, v becomes the root.
- u.p gets v as either u.left or u.right depending on whether u was u.p.left or u.p.right.
- Does not update v.left or v.right.

6.5. Deletion

- If z has no left child, replace by its right child. The right child may or may not be NIL (If NIL \rightarrow no child)




if ($z.\text{left} = \text{NIL}$) then
TRANSPLANT (T, z, r)




6.5. Deletion

- if z has just one child, the left child, then replace z by its left child.




if ($z.\text{right} = \text{NIL}$) then
TRANSPLANT (T, z, l)




6.5. Deletion



- if y is the right child of z , then replace z by y and leave the right child of y alone.




$y = \text{TREE-MINIMUM}(z.\text{right})$
if $(y.p = z)$ then
 $\text{TRANSPLANT}(T, z, y)$




6.5. Deletion

- Otherwise, y lie with in the right subtree of z , but it is not the root of this subtree. We replace y by its own right child. Then we replace z by y .




```

y = TREE-MINIMUM(z.right)
if (y.p ≠ z) then
    TRANSPLANT (T, y, y.right)
    y.right = z.right
    y.right.p = y
    TRANSPLANT(T, z, y)
    y.left = z.left
    y.left.p = y
    
```




6.5. Deletion

- Otherwise, y lie with in the right subtree of z , but it is not the root of this subtree. We replace y by its own right child. Then we replace z by y .






```

y = TREE-MINIMUM(z.right)
if (y.p ≠ z) then
    TRANSPLANT (T, y, y.right)
    y.right = z.right
    y.right.p = y
    TRANSPLANT(T, z, y)
    y.left = z.left
    y.left.p = y
    
```



6.5. Deletion

- Otherwise, y lies within the right subtree of z , but it is not the root of this subtree. We replace y by its own right child. Then we replace z by y .



6.5. Deletion




Algorithm (TREE-DELETE(T, z))

```
(1)   if (z.left = NIL) then           //no left child
(2)     TRANSPLANT( $T, z, z.right$ )
(3)   else
(4)     if (z.right = NIL) then        //no right child
(5)       TRANSPLANT( $T, z, z.left$ )
(6)     else                          //two children
(7)        $y = \text{TREE\_MINIMUM}(z.right)$ 
(8)       if ( $y.p \neq z$ ) then         //y is not z.right
(9)         TRANSPLANT( $T, y, y.right$ )
(10)       $y.right = z.right$ 
(11)       $y.right.p = y$ 
(12)      TRANSPLANT( $T, z, y$ )         //y is z.right
(13)       $y.left = z.left$ 
(14)       $y.left.p = y$ 
```


6.5. Deletion

- Tree-delete(T, I): no left child




6.5. Deletion

- Tree-delete(T,G): has left, but no right child




6.5. Deletion

- Tree-delete(T,K): two children, successor is right child



6.5. Deletion

- Tree-delete(T,B): successor is not the right child.





6.5. Deletion

Time: $O(h)$, on a tree of height h .

Minimizing running time

We've been analyzing running time in terms of h (the height of the binary search tree), instead of n (the number of nodes in the tree).

- Problem: Worst case for binary search tree is $\Theta(n)$ - no better than linked list.
- Solution: Guarantee small height (balanced tree) with $h = O(\lg n)$.
- Method: We restructure the tree. Querying works as before (no adjustments necessary). Insertion or deletion (changing the structure) may require some extra work.

Red-black trees [Next Week] are a special class of binary trees that avoids the worst-case behavior of $O(n)$ like “plain” binary search trees.




6.6. Expected height of a randomly built binary search tree

- Given a set of n distinct keys.
- Insert them in random order into an initially empty binary search tree.
- We assume that each of the $n!$ permutations is equally likely.
- Different from assuming that every binary search tree on n keys is equally likely.
- Try it for $n = 3$. Will get 5 different binary search trees. When we look at the binary search trees resulting from each of the $3!$ input permutations, 4 trees will appear once and 1 tree will appear twice.
- Forget about deleting keys.
- We will show that the expected height of a randomly built binary search tree is $O(\lg n)$.



6.6. Expected height of a randomly built binary search tree

Build all possible BST when $A=[1,2,3]$





6.6. Expected height of a randomly built binary search tree

Random variables

Define the following random variables:

X_n = height of a randomly built binary search tree on n keys. $Y_n = 2^{X_n}$ = *exponential height*.

R_n = rank of the root within the set of n keys used to build the binary search tree.

- Equally likely to be any element of $\{1, 2, \dots, n\}$.
- If $R_n = i$, then
 - Left subtree is a randomly-built binary search tree on $i - 1$ keys.
 - Right subtree is a randomly-built binary search tree on $n - i$ keys.




6.6. Expected height of a randomly built binary search tree

Foreshadowing


We will need to relate $E[Y_n]$ to $E[X_n]$.

We'll use *Jensen's inequality*: $E[f(X)] \geq f(E[X])$ provided

- the expectations exist and are finite, and
 $f(x)$ is *convex*: for all x, y and all $0 \leq \lambda \leq 1$
- $f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$.



- We'll use Jensen's inequality for $f(x) = 2^x$.
- Since 2^x curves upward, it's convex.




6.6. Expected height of a randomly built binary search tree

Formula for Y_n

Think about Y_n , if we know that $R_n = i$:

Height of root is 1 more than the maximum height of its children:

$$Y_n = 2 \cdot \max(Y_{i-1}, Y_{n-i}).$$





6.6. Expected height of a randomly built binary search tree

Formula for Y_n

Base cases:

- $Y_1 = 1$ (expected height of a 1-node tree is $2^0 = 1$).
- Define $Y_0 = 0$.

Define indicator random variables $Z_{n,1}, Z_{n,2}, \dots, Z_{n,n} : Z_{n,i} = I\{R_n = i\}$.

R_n is equally likely to be any element of $\{1, 2, \dots, n\}$

$$\Rightarrow \Pr\{R_n = i\} = 1/n$$

$$\Rightarrow E[Z_{n,i}] = 1/n \text{ (since } E[I\{A\}] = \Pr\{A\} \text{) [Remember this]}$$

Consider a given n -node binary search tree (which could be a subtree). Exactly one $Z_{n,i}$ is 1, and all others are 0. Hence,

$$Y_n = \sum_{i=1}^n Z_{n,i} \cdot (2 \cdot \max(Y_{i-1}, Y_{n-i})).$$



6.6. Expected height of a randomly built binary search tree

Bounding $E[Y_n]$

We will show that $E[Y_n]$ is polynomial in n , which will imply that $E[X_n] = O(\lg n)$.

Claim

$Z_{n,1}$ is independent of Y_{n-1} and Y_{n-i} .

Justification: If we choose the root such that $R_n = i$, the left subtree contains $i - 1$ nodes, and it's like any other randomly built binary search tree with $i - 1$ nodes. Other than the number of nodes, the left subtree's structure has nothing to do with it being the left subtree of the root.

Hence, $Y_i - 1$ and $Z_{n,i}$ are independent.

Similarly, Y_{i-1} and $Z_{n,i}$ are independent.

Fact

If X and Y are nonnegative random variables, then $E[\max(X, Y)] \leq E[X] + E[Y]$.



6.6. Expected height of a randomly built binary search tree

Fact

If X and Y are nonnegative random variables, then $E[\max(X, Y)] \leq E[X] + E[Y]$.

$$E[Y_n] \leq \frac{2}{n} \sum_{i=1}^n (E[Y_{i-1}] + E[Y_{n-i}])$$

$$E[Y_0] + E[Y_{n-1}] + E[Y_1] + E[Y_{n-2}] + \cdots + E[Y_{n-1}] + E[Y_0] = 2 \sum_{i=0}^{n-1} E[Y_i]$$

$$E[Y_n] \leq \frac{4}{n} \sum_{i=0}^{n-1} E[Y_i]$$



6.6. Expected height of a randomly built binary search tree

Solving the recurrence

We will show that for all integers $n > 0$, this recurrence has the solution

$$E[Y_n] \leq \frac{4}{n} \binom{n+3}{3}$$

Lemma:

$$\sum_{i=0}^{n-1} \binom{i+3}{3} = \binom{n+3}{4}$$

Using Pascal's identity:

$$\begin{aligned} \binom{n}{k} &= \binom{n-1}{k-1} + \binom{n-1}{k} \\ \binom{n+3}{4} &= \sum_{i=0}^{n-1} \binom{i+3}{3} \end{aligned}$$

We solve the recurrence by induction on n .



6.6. Expected height of a randomly built binary search tree

Basis: $n = 1$

$$1 = Y_1 = E[Y_1] \leq \frac{1}{4} \binom{1+3}{3} = \frac{4}{4} = 1$$

Inductive step: Assume that $E[Y_n] \leq \frac{4}{n} \binom{n+3}{3}$ for all $i < n$. Then

$$E[Y_n] \leq \frac{4}{n} \sum_{i=0}^{n-1} E[Y_i] = \frac{1}{4} \binom{n+3}{3}.$$



6.6. Expected height of a randomly built binary search tree

Bounding $E[X_n]$

With our bound on $E[Y_n]$, we use Jensen's inequality to bound

$$E[X_n]: 2^{E[X_n]} \leq E[2^{X_n}] = E[Y_n].$$

Thus,

$$2^{E[X_n]} \leq \frac{1}{4} \binom{n+3}{3} = \frac{1}{4} \cdot \frac{(n+3)(n+2)(n+1)}{6} = cn^3$$

Taking logs of both sides gives $E[X_n] = O(\lg n)$.