



CS 590 – Algorithms

Lecture 7 – Red-Black Trees



Outlines

7. Red-black trees (RBTs)

7.1. Characteristics and Properties

7.2. Black Heights

7.3. Operations

 7.3.1. Rotations

 7.3.2. Insertions

 7.3.3. Deletions



7. Red-black trees

Red-black trees

- A variation of binary search trees.
- *Balanced*: height is $O(\lg n)$, where n is the number of nodes.
- Operations will take $O(\lg n)$ time in the worst case.



7.1. Properties

Red-black trees:

- A **red-black tree** is a self-balancing binary search tree with one extra (+1) bit per node:
 - An attribute *color*, which is either **red** or **black**.
 - All leaves are empty (*nil*) and colored **black**.
- We use a single sentinel, $T.nil$, for all the leaves of red-black tree $T.nil.color$ is **black**.
- The root's parent is also $T.nil$.
- All other attributes of binary search trees are inherited by red-black trees (*key*, *left*, *right*, and *p*).



7.1. Properties

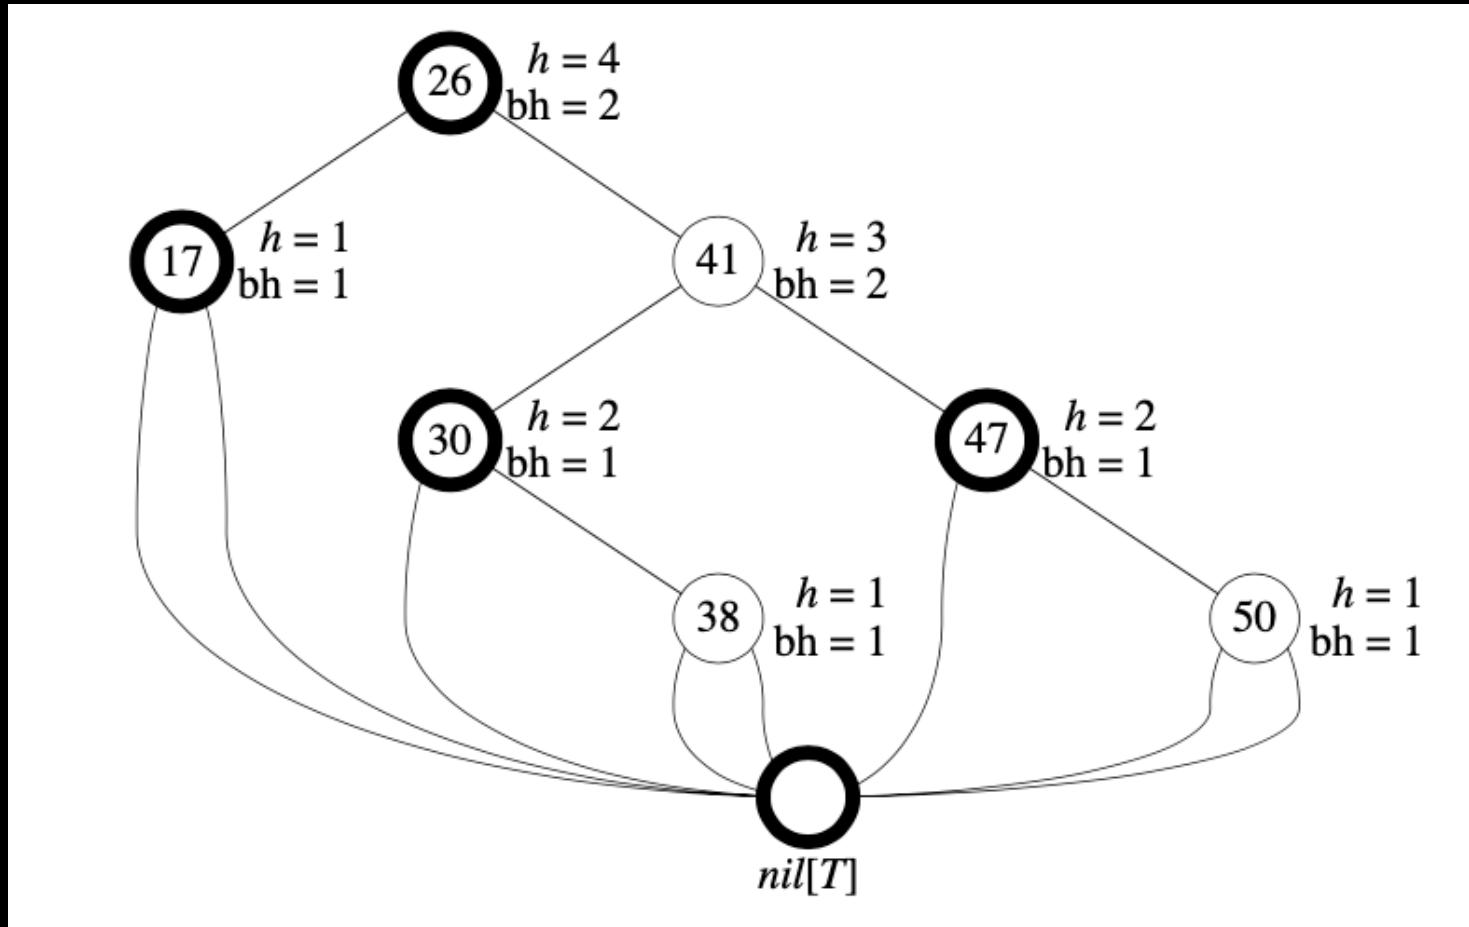
Red-black trees properties:

1. Every node is either red or black.
2. The root is black.
3. Every leaf ($T.nil$) is black.
4. If a node is red, then **both its children are black**. (Hence no two reds in a row on a simple path from the root to a leaf.)
5. For each node, all paths from the node to descendant leaves contain **the same number** of black nodes.

Height of a red-black tree:

- ***Height of a node*** is the number of edges in a longest path to a leaf.
- ***Black-height*** of a node x : $bh(x)$ is **the number of black nodes** (including $T.nil$) on the path from x to leaf, **not counting x** . By property 5, black-height is well defined.

7.2. Black Heights





7.2. Black Heights

Claim

Any node with height h has black-height $\geq h/2$.

Claim

The subtree rooted at any node x contains $\geq 2^{bh(x)} - 1$ internal nodes.

Lemma

A red-black tree with n internal nodes has height $\leq 2 \lg(n + 1)$.



7.3. Operations

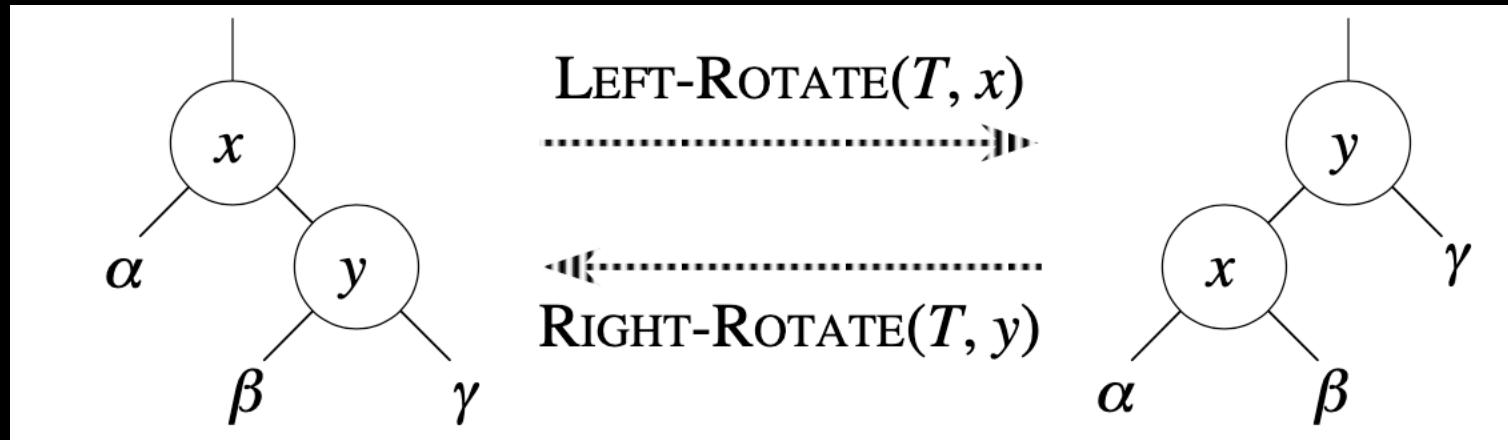
Operations on red-black trees

- The non-modifying binary-search-tree operations are unchanged.
- MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR, and SEARCH run in $O(h) = O(\lg n)$ time on red-black trees.
- Insertion and deletion are not so easy.
 - If we insert, what color to make the new node?
 - **Red:** Might violate property 4.
 - **Black:** Might violate property 5.
 - If we delete, thus removing a node, what color was the node that was removed?
 - **Red:** OK, since we won't have changed any black-heights, nor will we have created two red nodes in a row. Does not violate property 4.
 - Cannot violate property 2 – root is black.
 - **Black:** Could cause two reds in a row (violating property 4), and
 - can cause a violation of property 5.
 - Possibly violate property 2, if the removed node was the root and its child was red.

7.3.1. Rotations

Rotations:

- The basic tree-restructuring operation.
- Used to maintain red-black trees as balanced binary search trees.
- Changes the local pointer structure. (Only pointers are changed.)
- Won't upset the binary-search-tree property.
- Have both left rotation and right rotation. They are inverses of each other.
- A rotation takes a red-black-tree and a node within the tree.





7.3.1. Rotations

Algorithm (LEFT-ROTATE(T, x))

- (1) $y = x.\text{right}$ //set y
- (2) $x.\text{right} = y.\text{left}$ //turn left subtree of y into right of x
- (3) if ($y.\text{left} \neq T.\text{nil}$) then
- (4) $y.\text{left}.p = x$
- (5) $y.p = x.p$ //link parent of x to y
- (6) if ($x.p = T.\text{nil}$) then
- (7) $T.\text{root} = y$
- (8) else
- (9) if ($x = x.p.\text{left}$) then
- (10) $x.p.\text{left} = y$
- (11) else
- (12) $x.p.\text{right} = y$
- (13) $y.\text{left} = x$ //put x of left of y
- (14) $x.p = y$



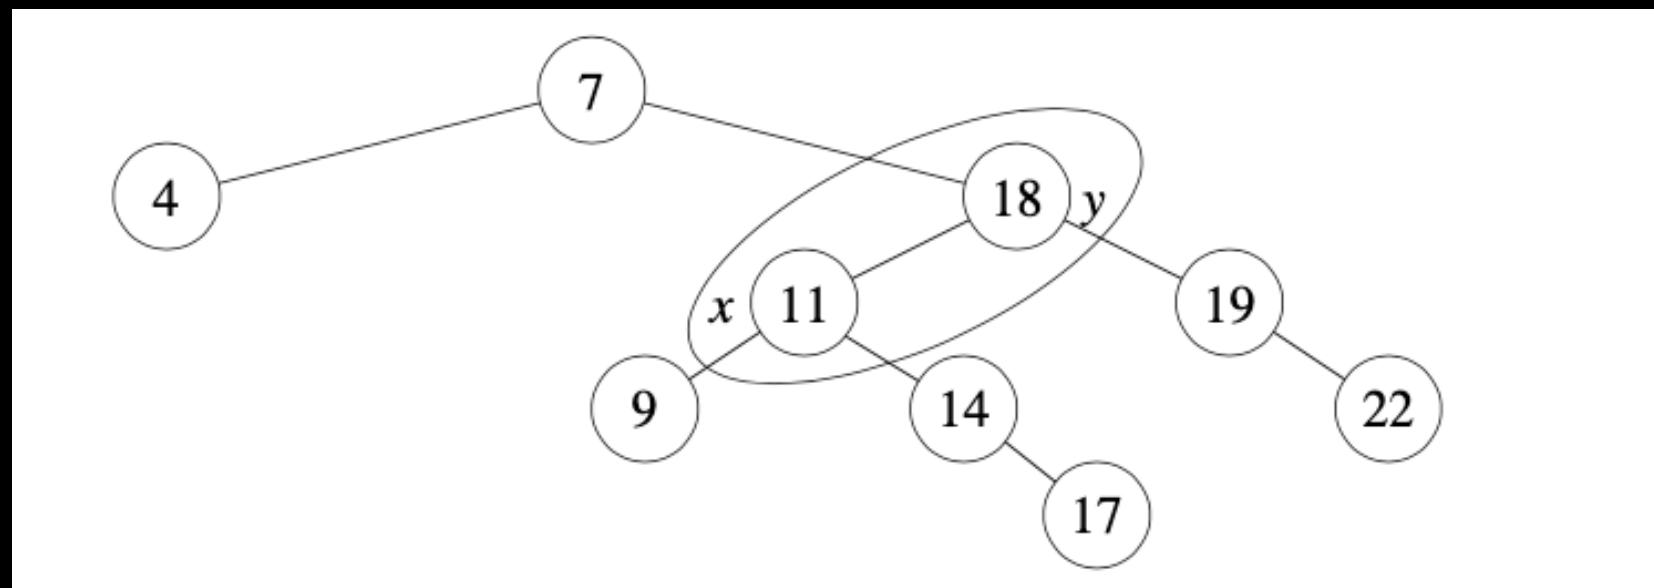
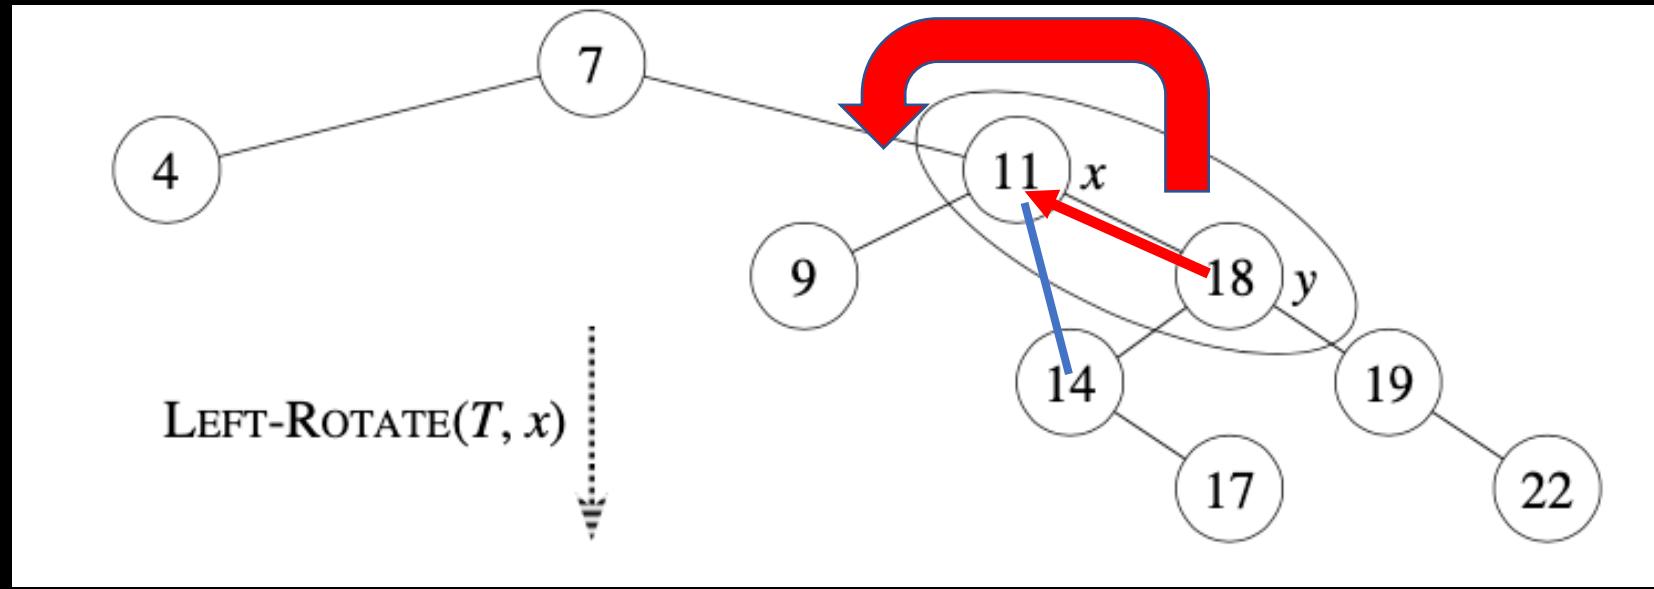
7.3.1. Rotations

The pseudocode for LEFT-ROTATE assumes that

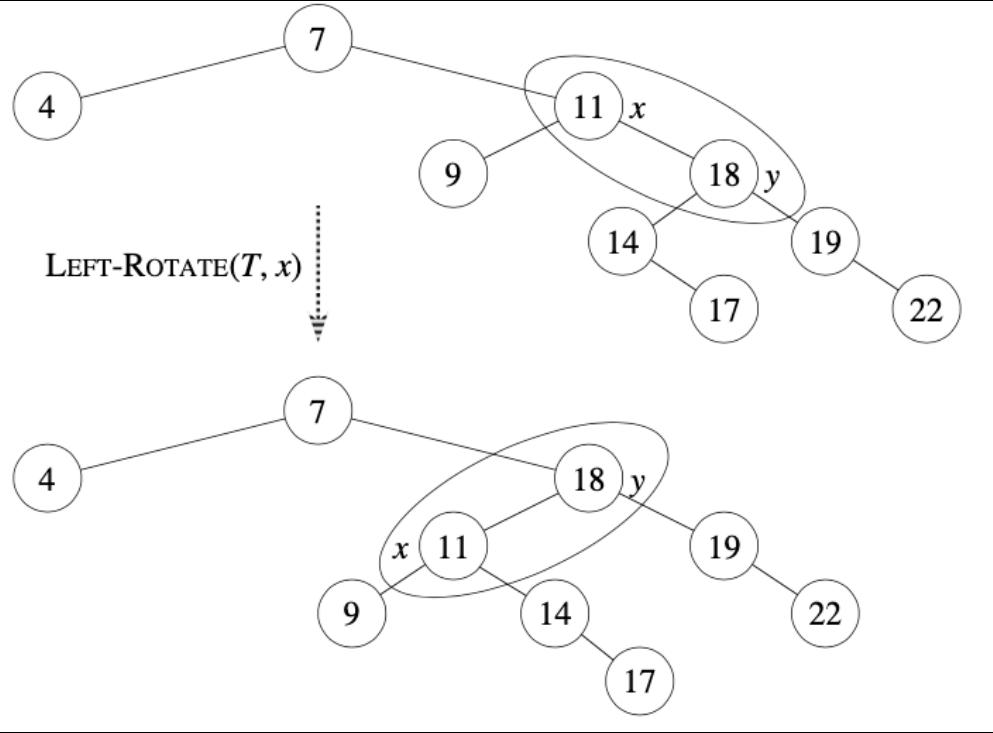
- $x.right \neq T.nil$, and
- root's parent is $T.nil$.

Pseudocode for RIGHT-ROTATE is symmetric: exchange *left* and *right* everywhere.

7.3.1. Rotations



7.3.1. Rotations



- Before rotation: keys of x 's left subtree $\leq 11 \leq$ keys of y 's left subtree $\leq 18 \leq$ keys of y 's right subtree.
- Rotation makes y 's left subtree into x 's right subtree.
- After rotation: keys of x 's left subtree $\leq 11 \leq$ keys of x 's right subtree $\leq 18 \leq$ keys of y 's right subtree.

Time: $O(1)$ for both LEFT-ROTATE and RIGHT-ROTATE, since a constant number of pointers are modified.

Notes:

- Rotation is a very basic operation, also used in Adelson-Velskii, Landis (AVL – BST whose h of right and left subtrees are differ by 1) trees and splay trees (BST w/ no explicit balance condition).
- Some books talk of rotating on an edge rather than on a node.



7.3.2. Insertions

Algorithm (RB-INSERT(T, x))

```
(1)       $y = T.\text{nil}$ ,  $x = T.\text{root}$ 
(2)      while ( $x \neq T.\text{nil}$ ) do
(3)           $y = x$ 
(4)          if ( $z.\text{key} < x.\text{key}$ ) then
(5)               $x = x.\text{left}$ 
(6)          else  $x = x.\text{right}$ 
(7)           $z.p = y$ 
(8)          if ( $y = NIL$ ) then
(9)               $T.\text{root} = z$ 
(10)         else if ( $z.\text{key} < y.\text{key}$ ) then
(11)              $y.\text{left} = z$ 
(12)         else  $y.\text{right} = z$ 
(13)          $z.\text{left} = T.\text{nil}$ 
(14)          $z.\text{right} = T.\text{nil}$ 
(15)          $z.\text{color} = \text{RED}$ 
(16)         RB-INSERT-FIXUP( $T, z$ )
```

7.3.2. Insertions

Algorithm (TREE-INSERT(T, x))

```

(1)      y = NIL, x = T.root
(2)      while ( $x \neq NIL$ ) do
(3)          y = x
(4)          if ( $z.key < x.key$ ) then
(5)              x = x.left
(6)          else x = x.right
(7)          z.p = y
(8)          if ( $y = NIL$ ) then
(9)              T.root = z
(10)         else if ( $z.key < y.key$ ) then
(11)             y.left = z
(12)         else y.right = z
    
```

Algorithm (RB-INSERT(T, x))

```

(1)      y = T.nil, x = T.root
(2)      while ( $x \neq T.nil$ ) do
(3)          y = x
(4)          if ( $z.key < x.key$ ) then
(5)              x = x.left
(6)          else x = x.right
(7)          z.p = y
(8)          if ( $y = NIL$ ) then
(9)              T.root = z
(10)         else if ( $z.key < y.key$ ) then
(11)             y.left = z
(12)         else y.right = z
(13)         z.left = T.nil
(14)         z.right = T.nil
(15)         z.color = RED
(16)         RB-INSERT-FIXUP( $T, z$ )
    
```



7.3.2. Insertions

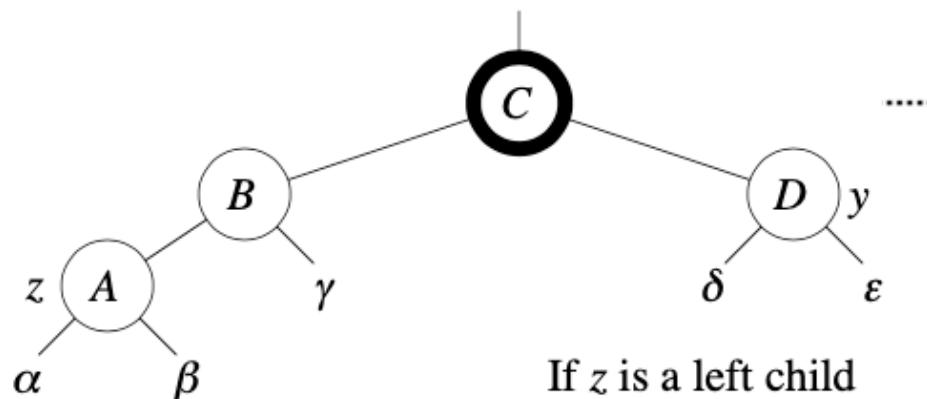
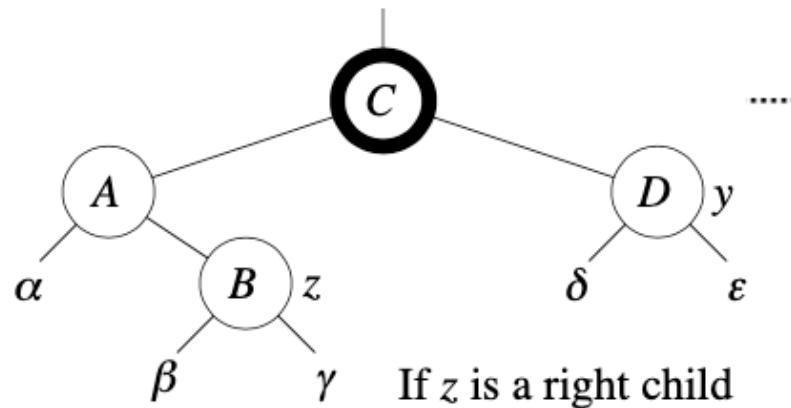
- RB-INSERT ends by coloring the new node z red.
- Then it calls RB-INSERT-FIXUP because we could have violated a red-black property.
- Which property might be violated?
 1. OK.
 2. If z is the root, then there's a violation. Otherwise, OK.
 3. OK.
 4. If $z.p$ is red, there's a violation: both z and $z.p$ are red.
 5. OK.

⇒ Remove the violation by calling RB-INSERT-FIXUP:

7.3.2. Insertions

when both z.p and y are red

Case 1: y is red

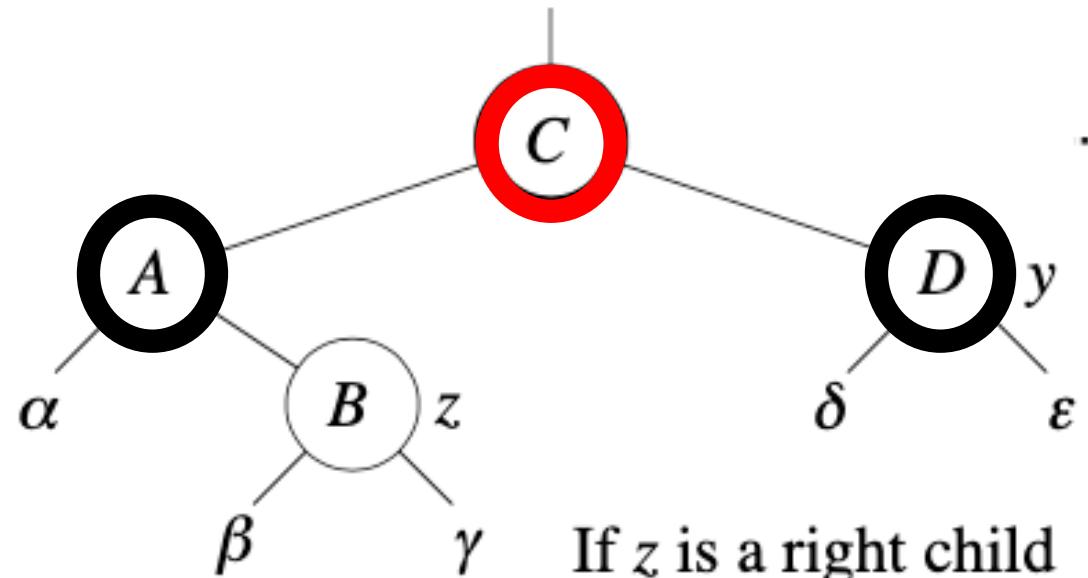


- **$z.p.p$ (z's grandparent) must be black**, since z and $z.p$ are both red and there are no other violations of property 4.
- **Make $z.p$ and y black** \Rightarrow now z and $z.p$ are not both red.
- But property 5 might now be violated.
- **Make $z.p.p$ red** \Rightarrow restores property 5.
- The next iteration has **$z.p.p$ as the new z** (i.e., z moves up 2 levels).

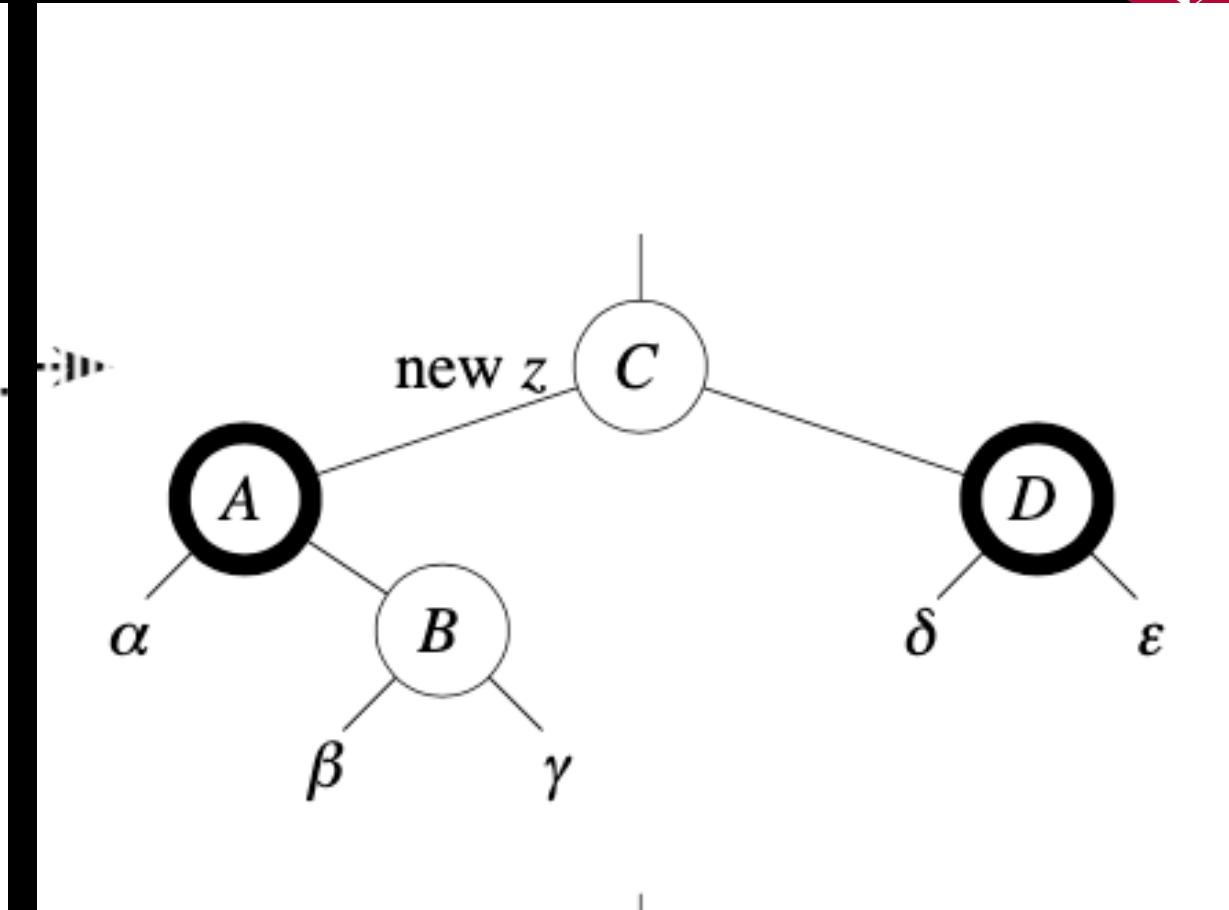
7.3.2. Insertions



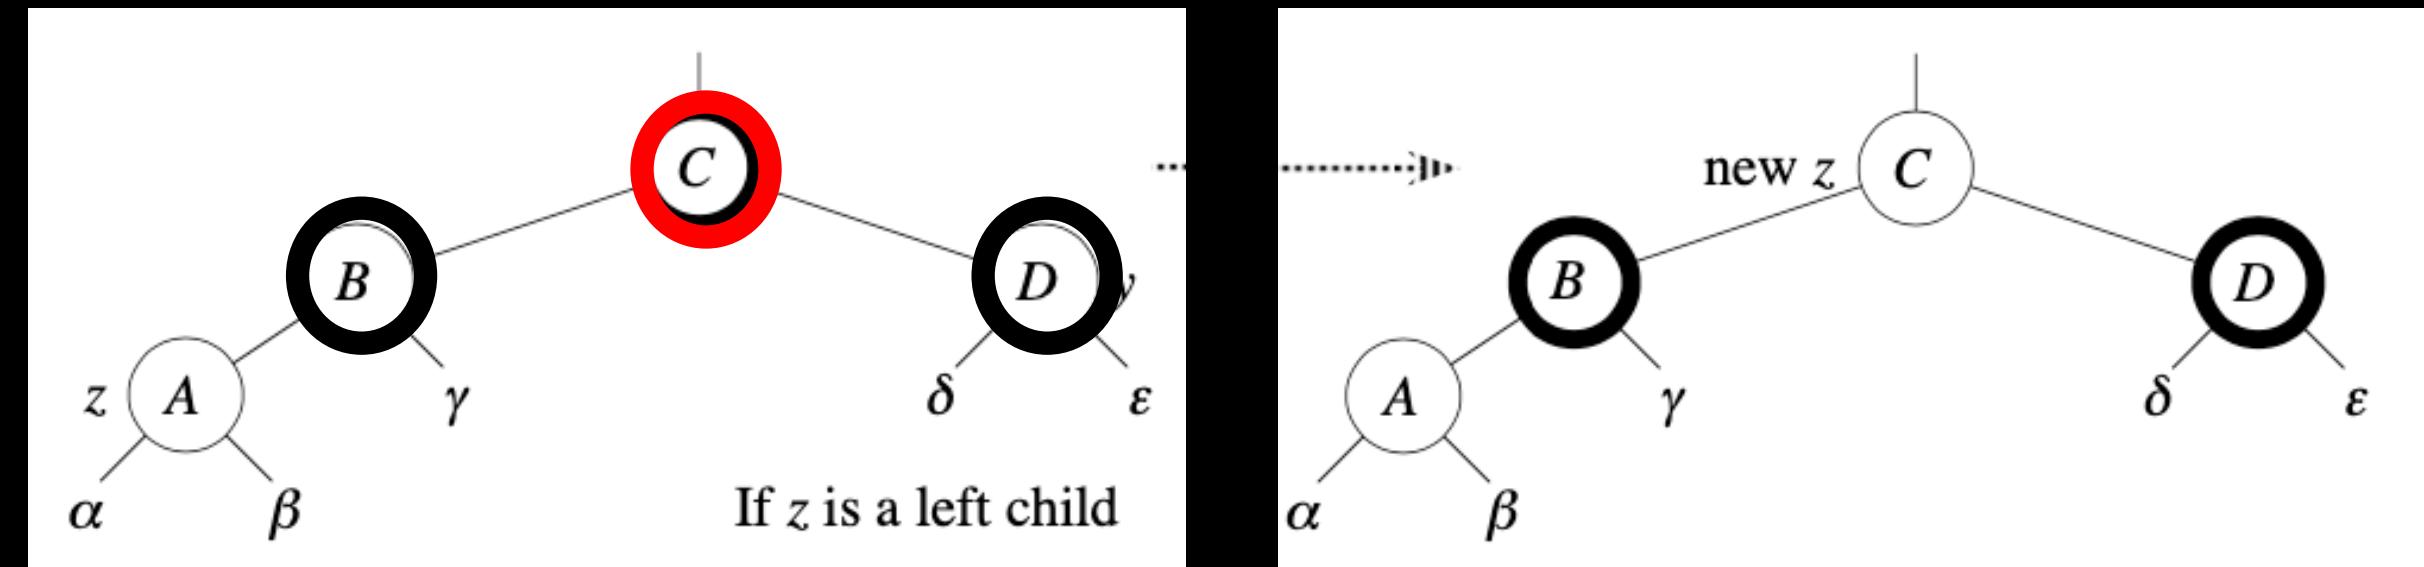
Case 1: y is red



If z is a right child

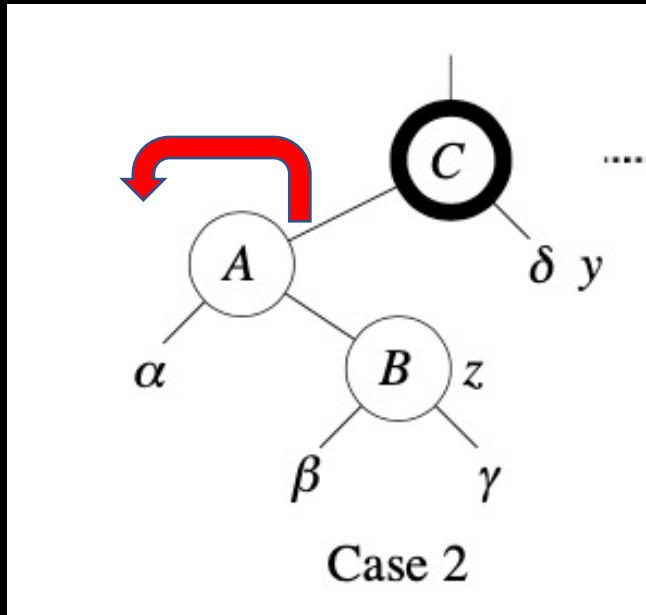


7.3.2. Insertions

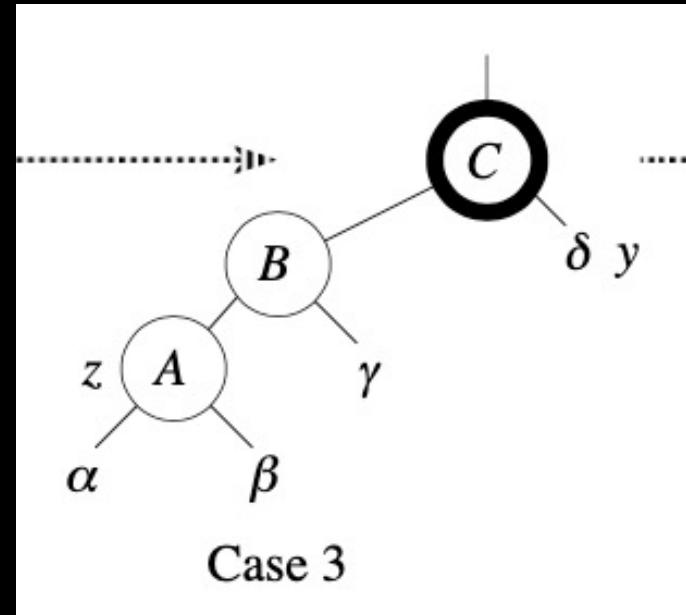


7.3.2. Insertions

Case 2: y is black, z is a right child

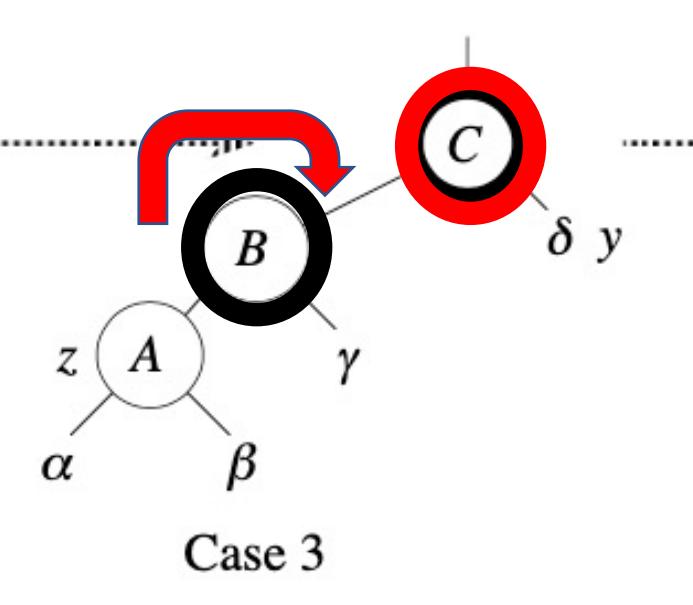


- **Left rotate around $z.p$** \Rightarrow now z is a left child, and **both z and $z.p$ are red**.
- Takes us immediately to case 3.

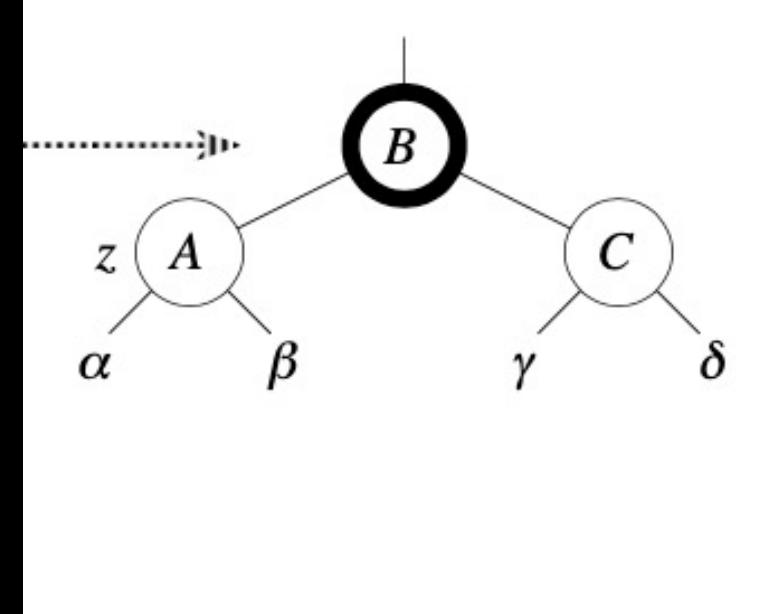


7.3.2. Insertions

Case 3: y is black, z is a left child



- **Make $z.p$ black and $z.p.p$ red.**
- Then **right rotate on $z.p.p$.**
- No longer have 2 reds in a row.
- **$z.p$ is now black** \Rightarrow no more iterations.





7.3.2. Insertions

Algorithm (RB-INSERT-FIXUP(T, z))

```
(1)      while ( $z.p.color = \text{RED}$ ) do
(2)          if ( $z.p = z.p.p.left$ ) then
(3)               $y = z.p.p.left$ 
(4)              if ( $y.color = \text{RED}$ ) then
(5)                   $z.p.color = \text{BLACK}$ ,  $y.color = \text{BLACK}$  //case 1
(6)                   $z.p.p.color = \text{RED}$ 
(7)                   $z = z.p.p$ 
(8)              else
(9)                  if ( $z = z.p.right$ ) then
(10)                      $z = z.p$                                 //case 2
(11)                     LEFT-ROTATE( $T, z$ )
(12)                      $z.p.color = \text{BLACK}$ ,  $z.p.p.color = \text{RED}$  //case 3
(13)                     RIGHT-ROTATE( $T, z.p.p$ )
(14)                 else
(15)                     (do same with right and left exchange)
(16)              $T.root.color = \text{BLACK}$ 
```

7.3.2. Insertions

Algorithm (RB-INSERT-FIXUP(T, z))

```

(1)      while ( $z.p.color = \text{RED}$ ) do
(2)          if ( $z.p = z.p.p.left$ ) then
(3)               $y = z.p.p.left$ 
(4)              if ( $y.color = \text{RED}$ ) then
(5)                   $z.p.color = \text{BLACK}$ ,  $y.color = \text{BLACK}$ 
(6)                   $z.p.p.color = \text{RED}$ 
(7)                   $z = z.p.p$ 
(8)              else
(9)                  if ( $z = z.p.right$ ) then
(10)                      $z = z.p$                                 //case 2
(11)                     LEFT-ROTATE( $T, z$ )
(12)                      $z.p.color = \text{BLACK}$ ,  $z.p.p.color = \text{RED}$  //case 3
(13)                     RIGHT-ROTATE( $T, z.p.p$ )
(14)                 else
(15)                     (do same with right and left exchange)
(16)              $T.root.color = \text{BLACK}$ 

```

Loop invariant:

At the start of each iteration of the **while** loop,

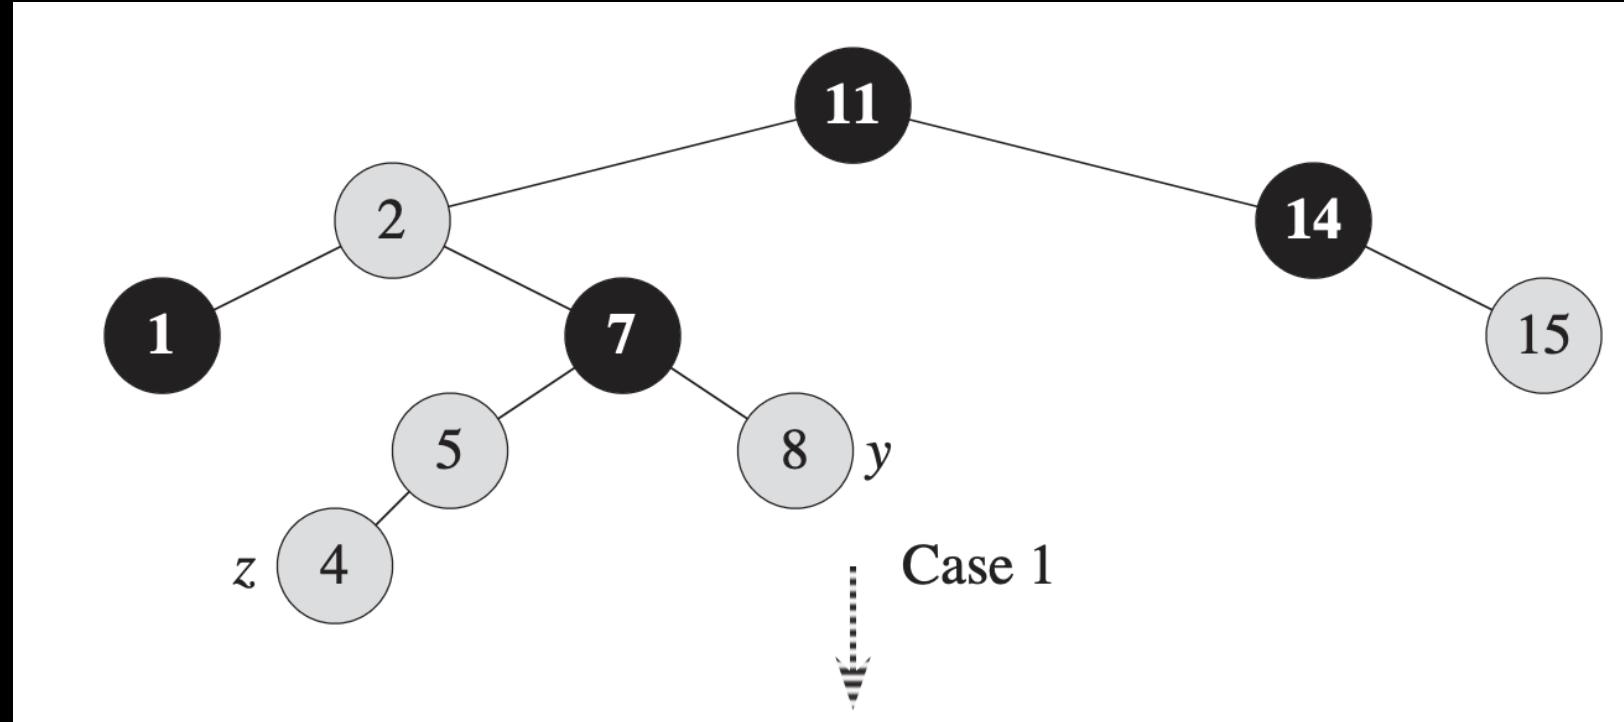
1. z is red.
2. If $z.p$ is the root, then $z.p$ is black.
3. There is at most one red-black violation:
 - a. Property 2: z is a red root, or
 - b. Property 4: z and $z.p$ are both red.



7.3.2. Insertions

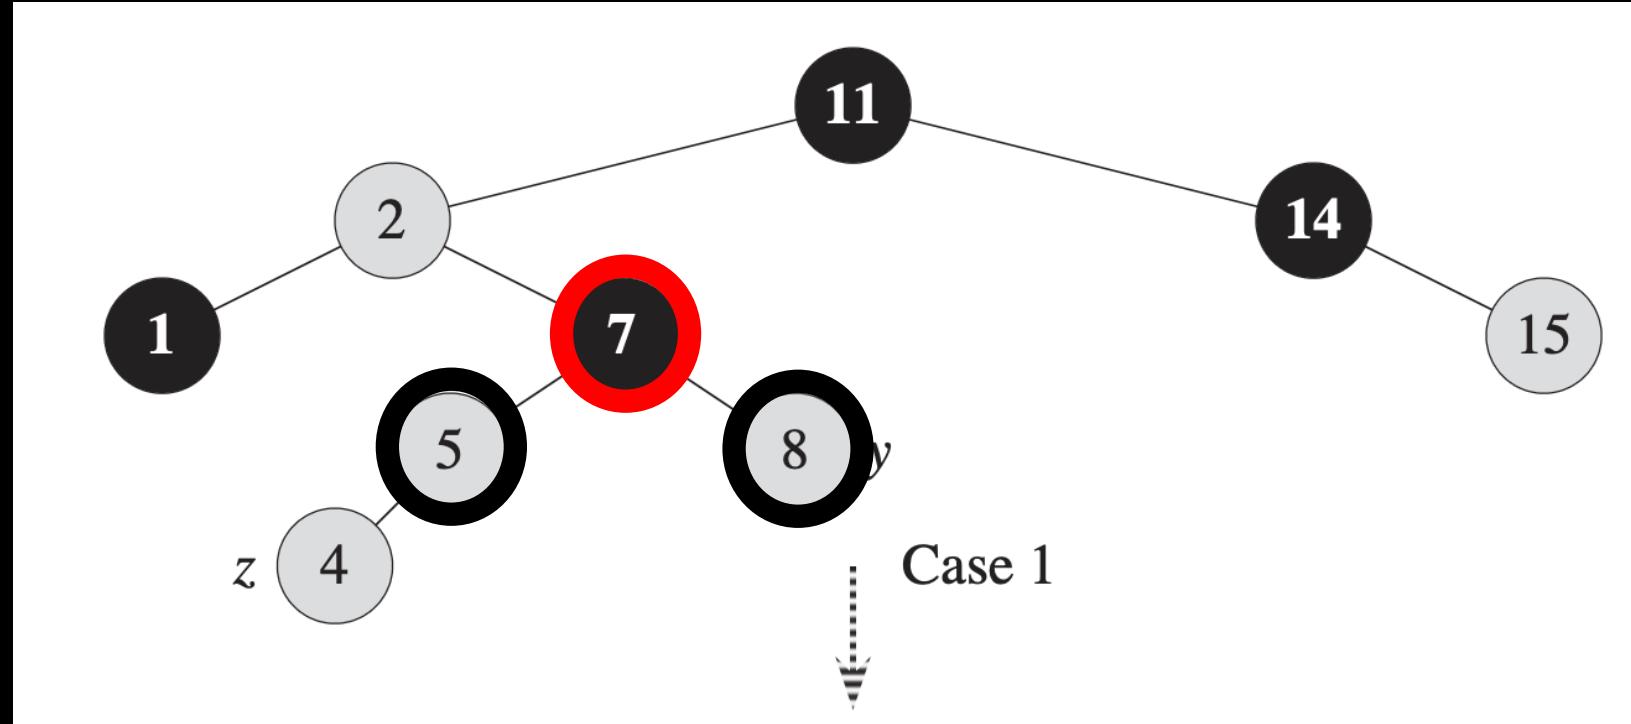
- **Initialization:** We've already seen why the loop invariant holds initially.
- **Termination:** The loop terminates because $z.p$ is black. Hence, property 4 is OK. Only property 2 might be violated, and the last line fixes it.
- **Maintenance:** We drop out when z is the root (since then $z.p$ is the sentinel $T.nil$, which is black). When we start the loop body, the only violation is of property 4.
- There are 6 cases, 3 of which are symmetric to the other 3. The cases are not mutually exclusive. We'll consider cases in which $z.p$ is a left child.
- Let y be z 's uncle ($z.p$'s sibling).

7.3.2. Insertions

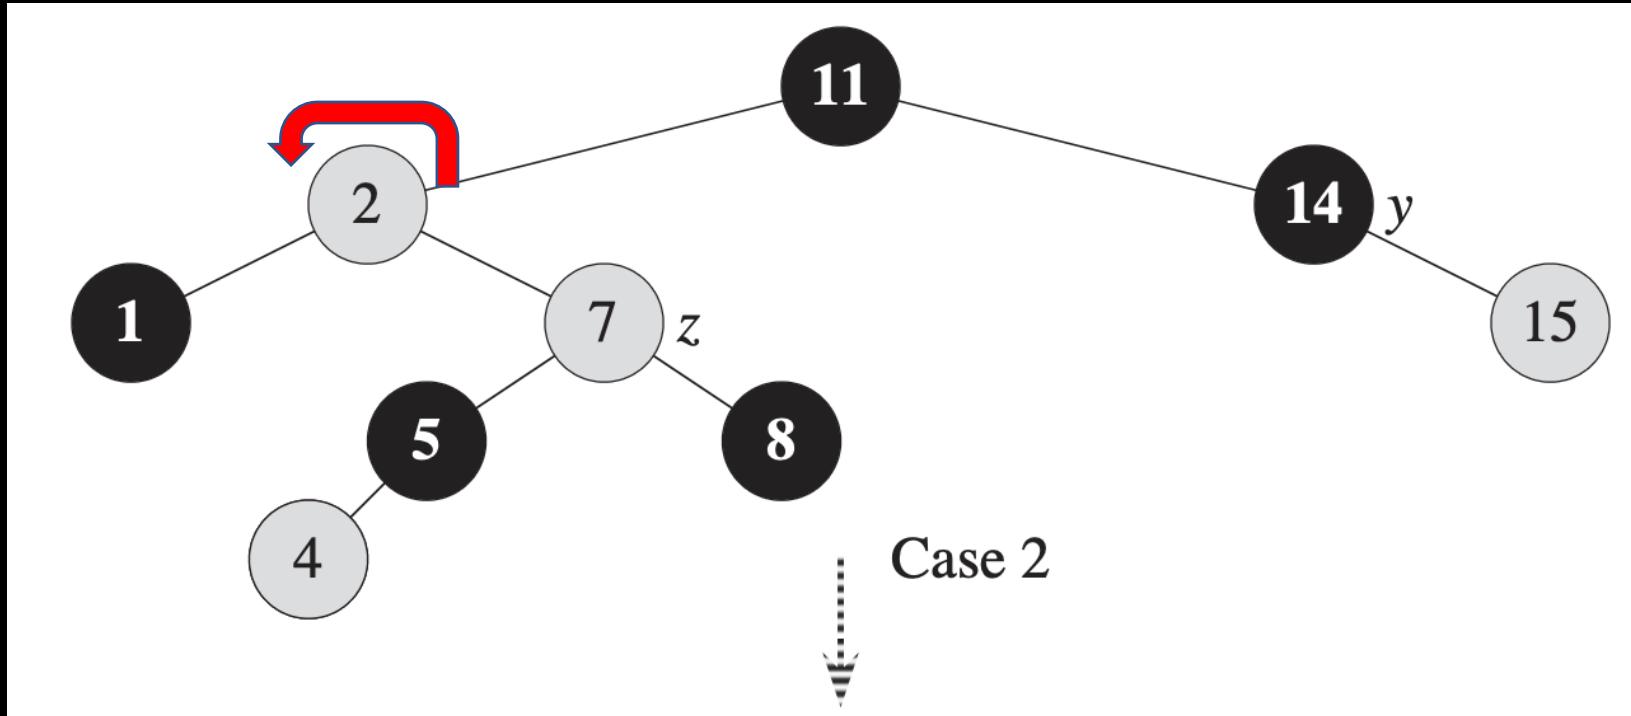


- Which of the red-black properties might be violated upon the call to RB-INSERT-FIXUP?
- Property 1, 3, 5 are satisfied.
- Property 2 and property 4 possible violations are due to z being colored red.
- Property 2 is violated if z is the root, and property 4 is violated if z's parent is red.

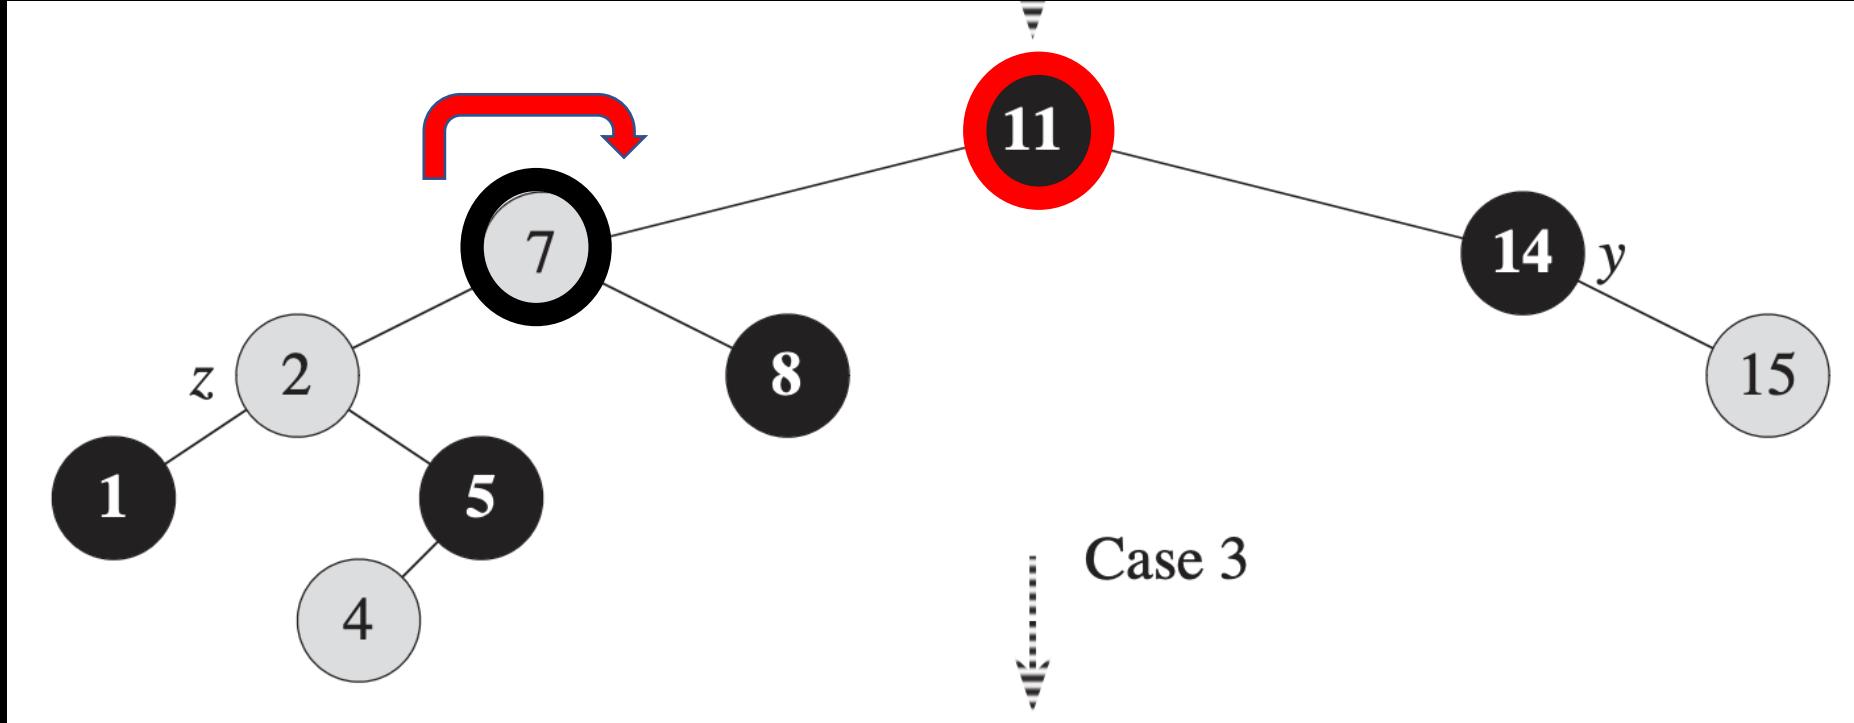
7.3.2. Insertions



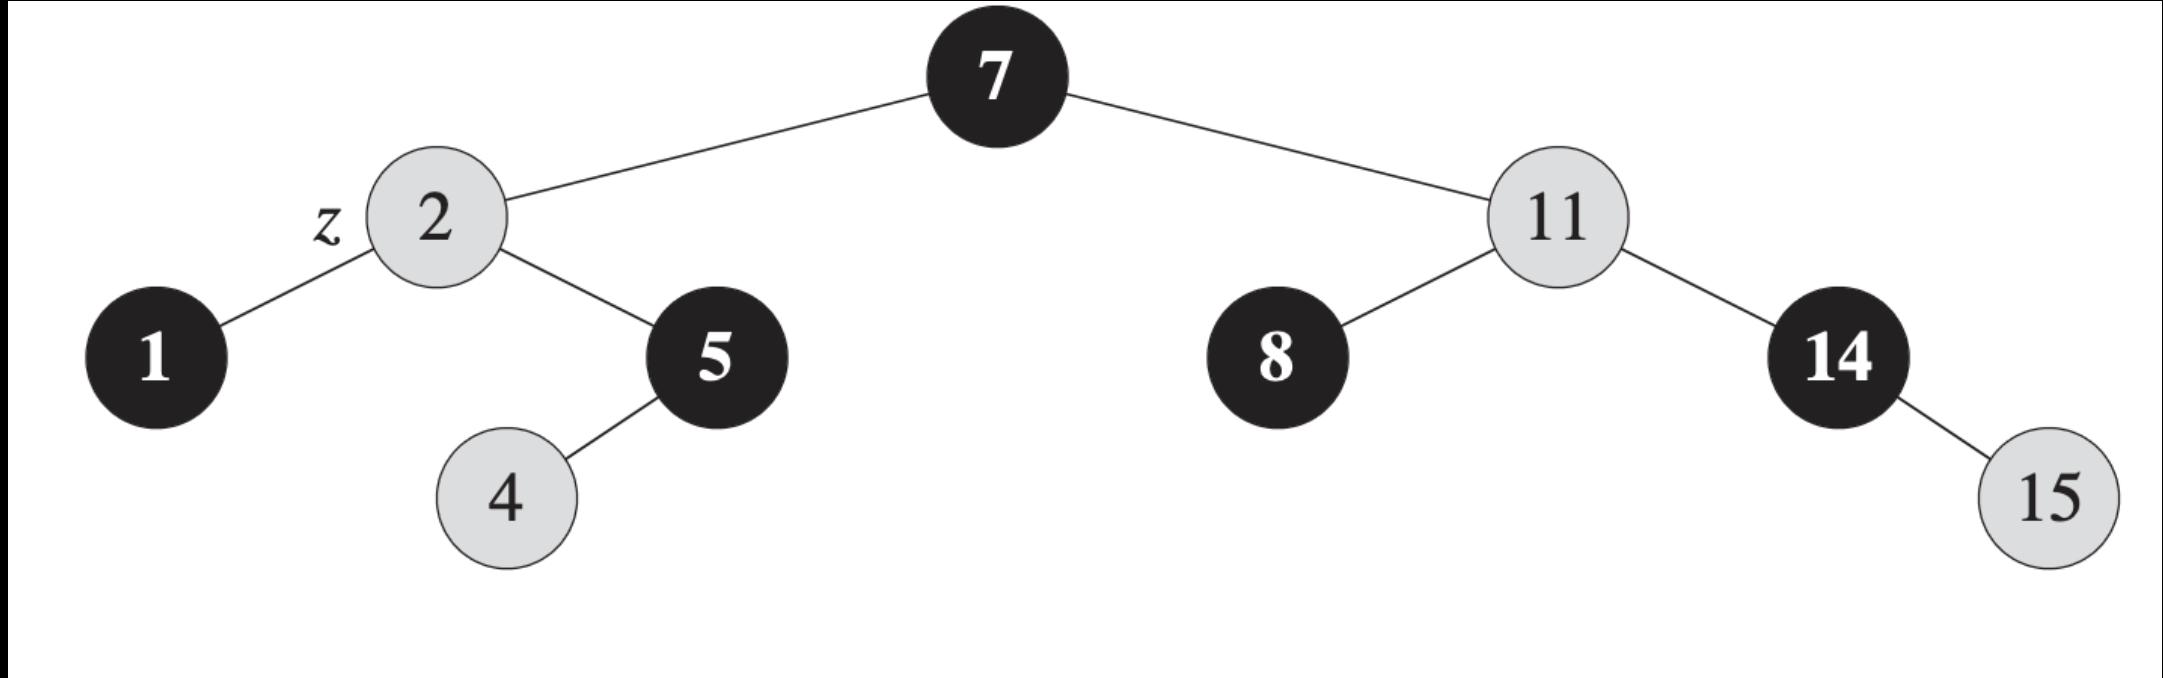
7.3.2. Insertions



7.3.2. Insertions



7.3.2. Insertions





7.3.2. Insertions

Analysis

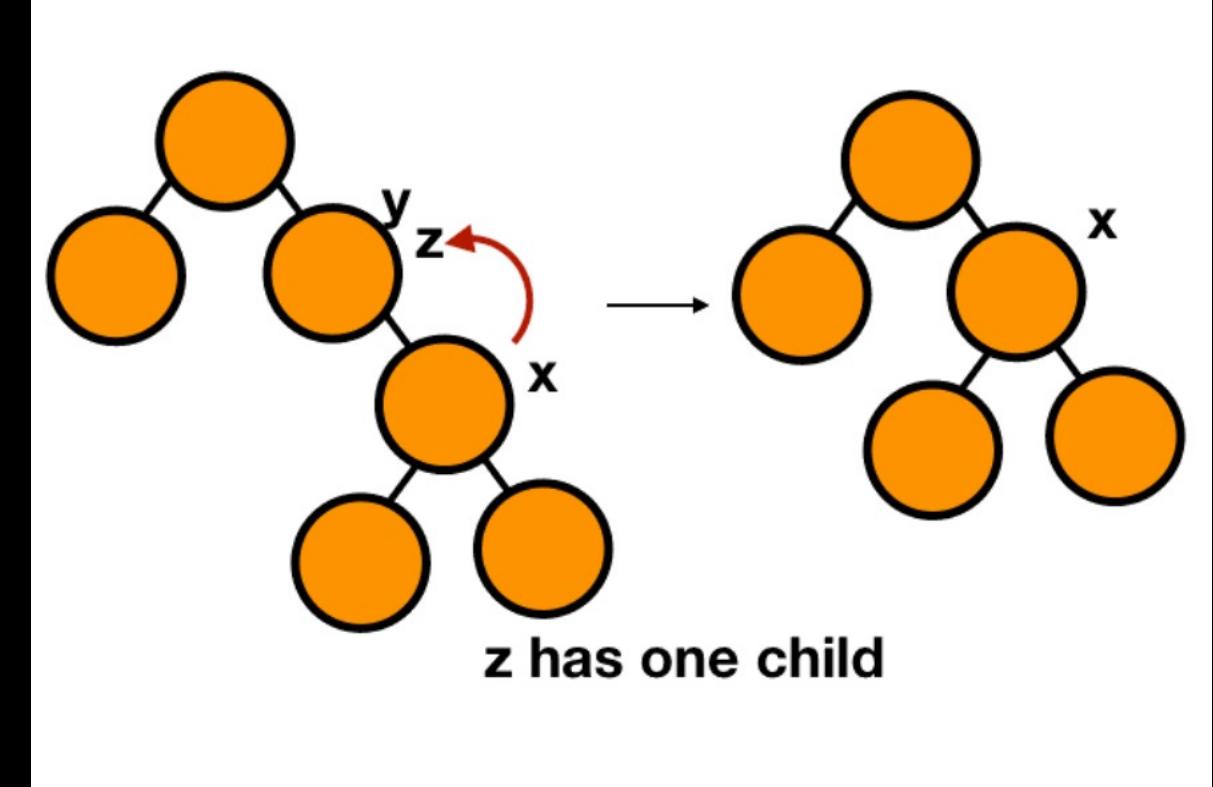
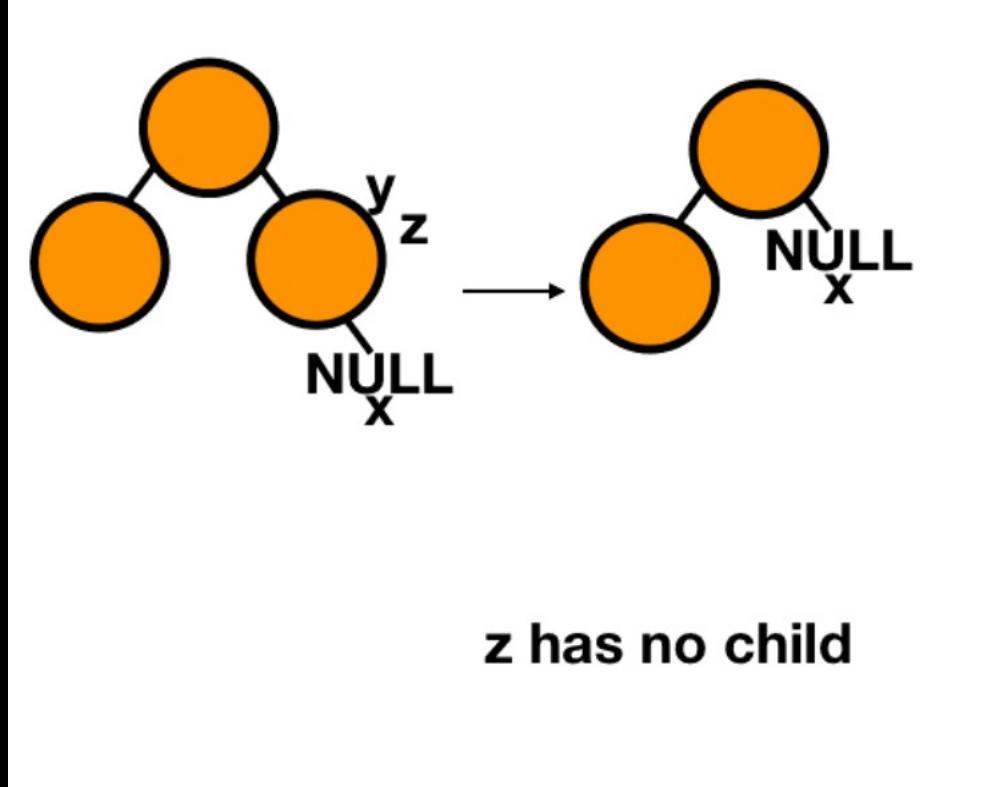
$O(\lg n)$ time to get through RB-INSERT up to the call of RB-INSERT-FIXUP.

Within RB-INSERT-FIXUP:

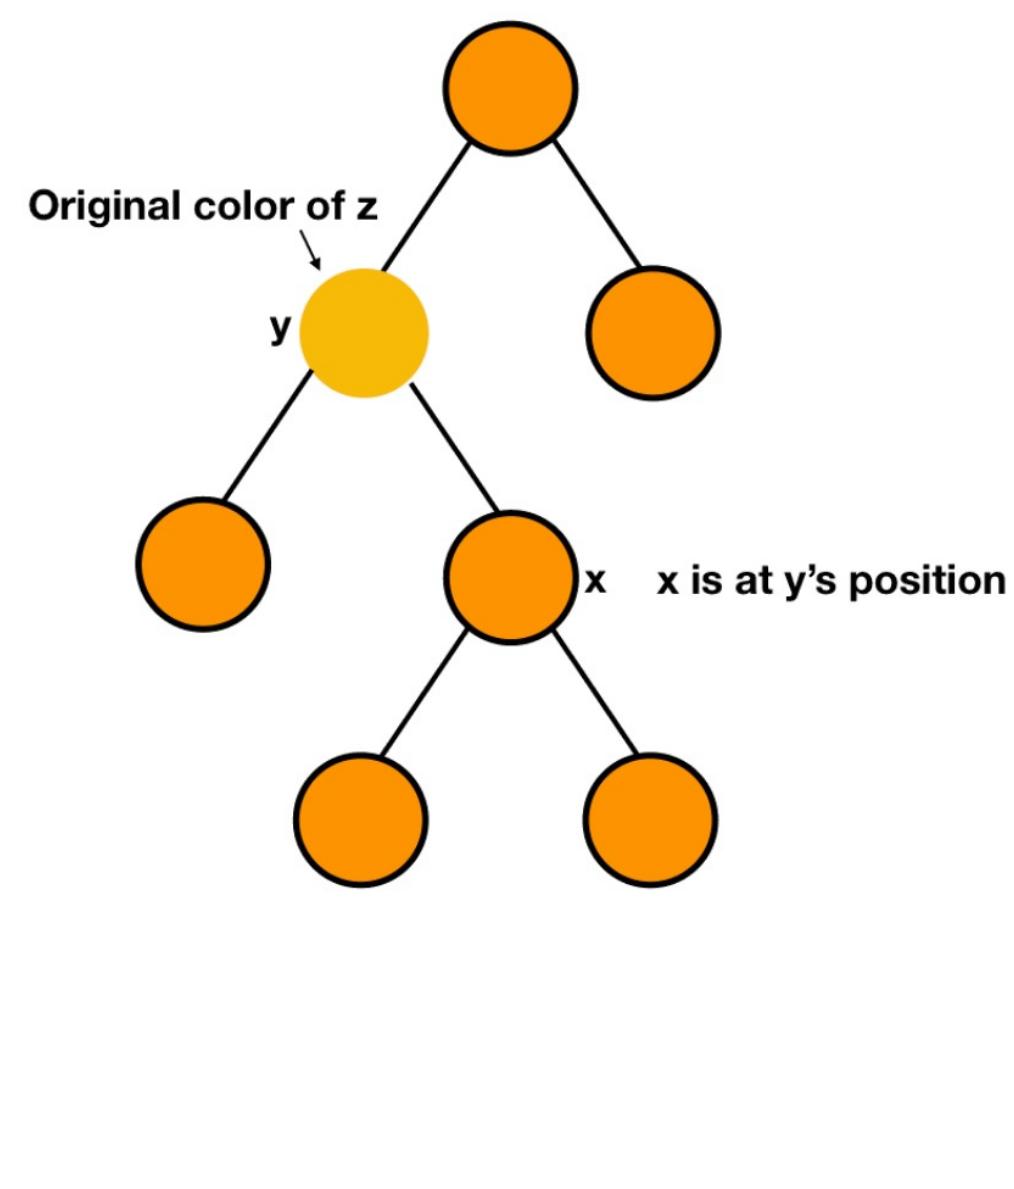
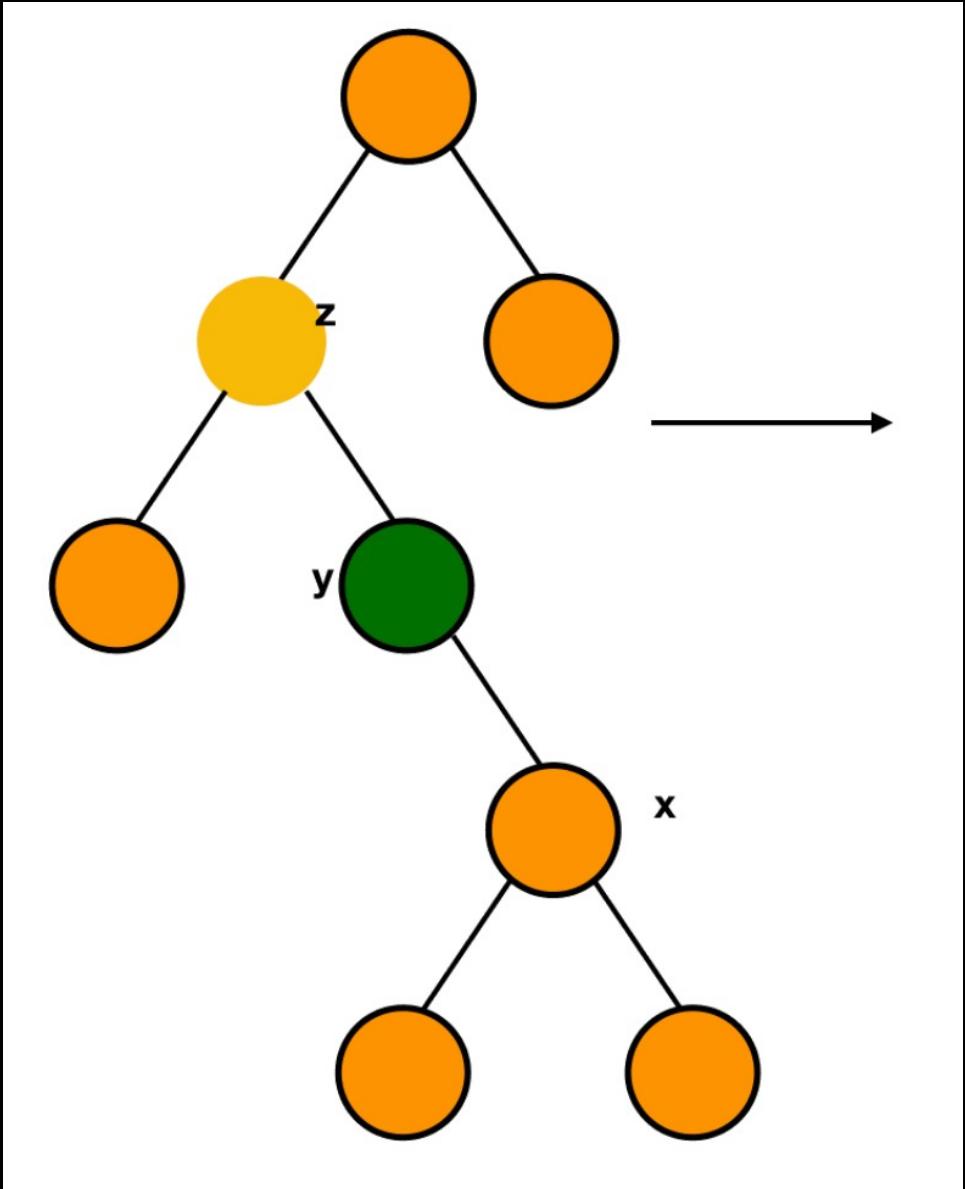
- Each iteration takes $O(1)$ time.
- Each iteration is either the last one or it moves z up 2 levels.
- $O(\lg n)$ levels $\Rightarrow O(\lg n)$ time.
- Also note that there are at most 2 rotations overall.

Thus, insertion into a red-black tree takes $O(\lg n)$ time.

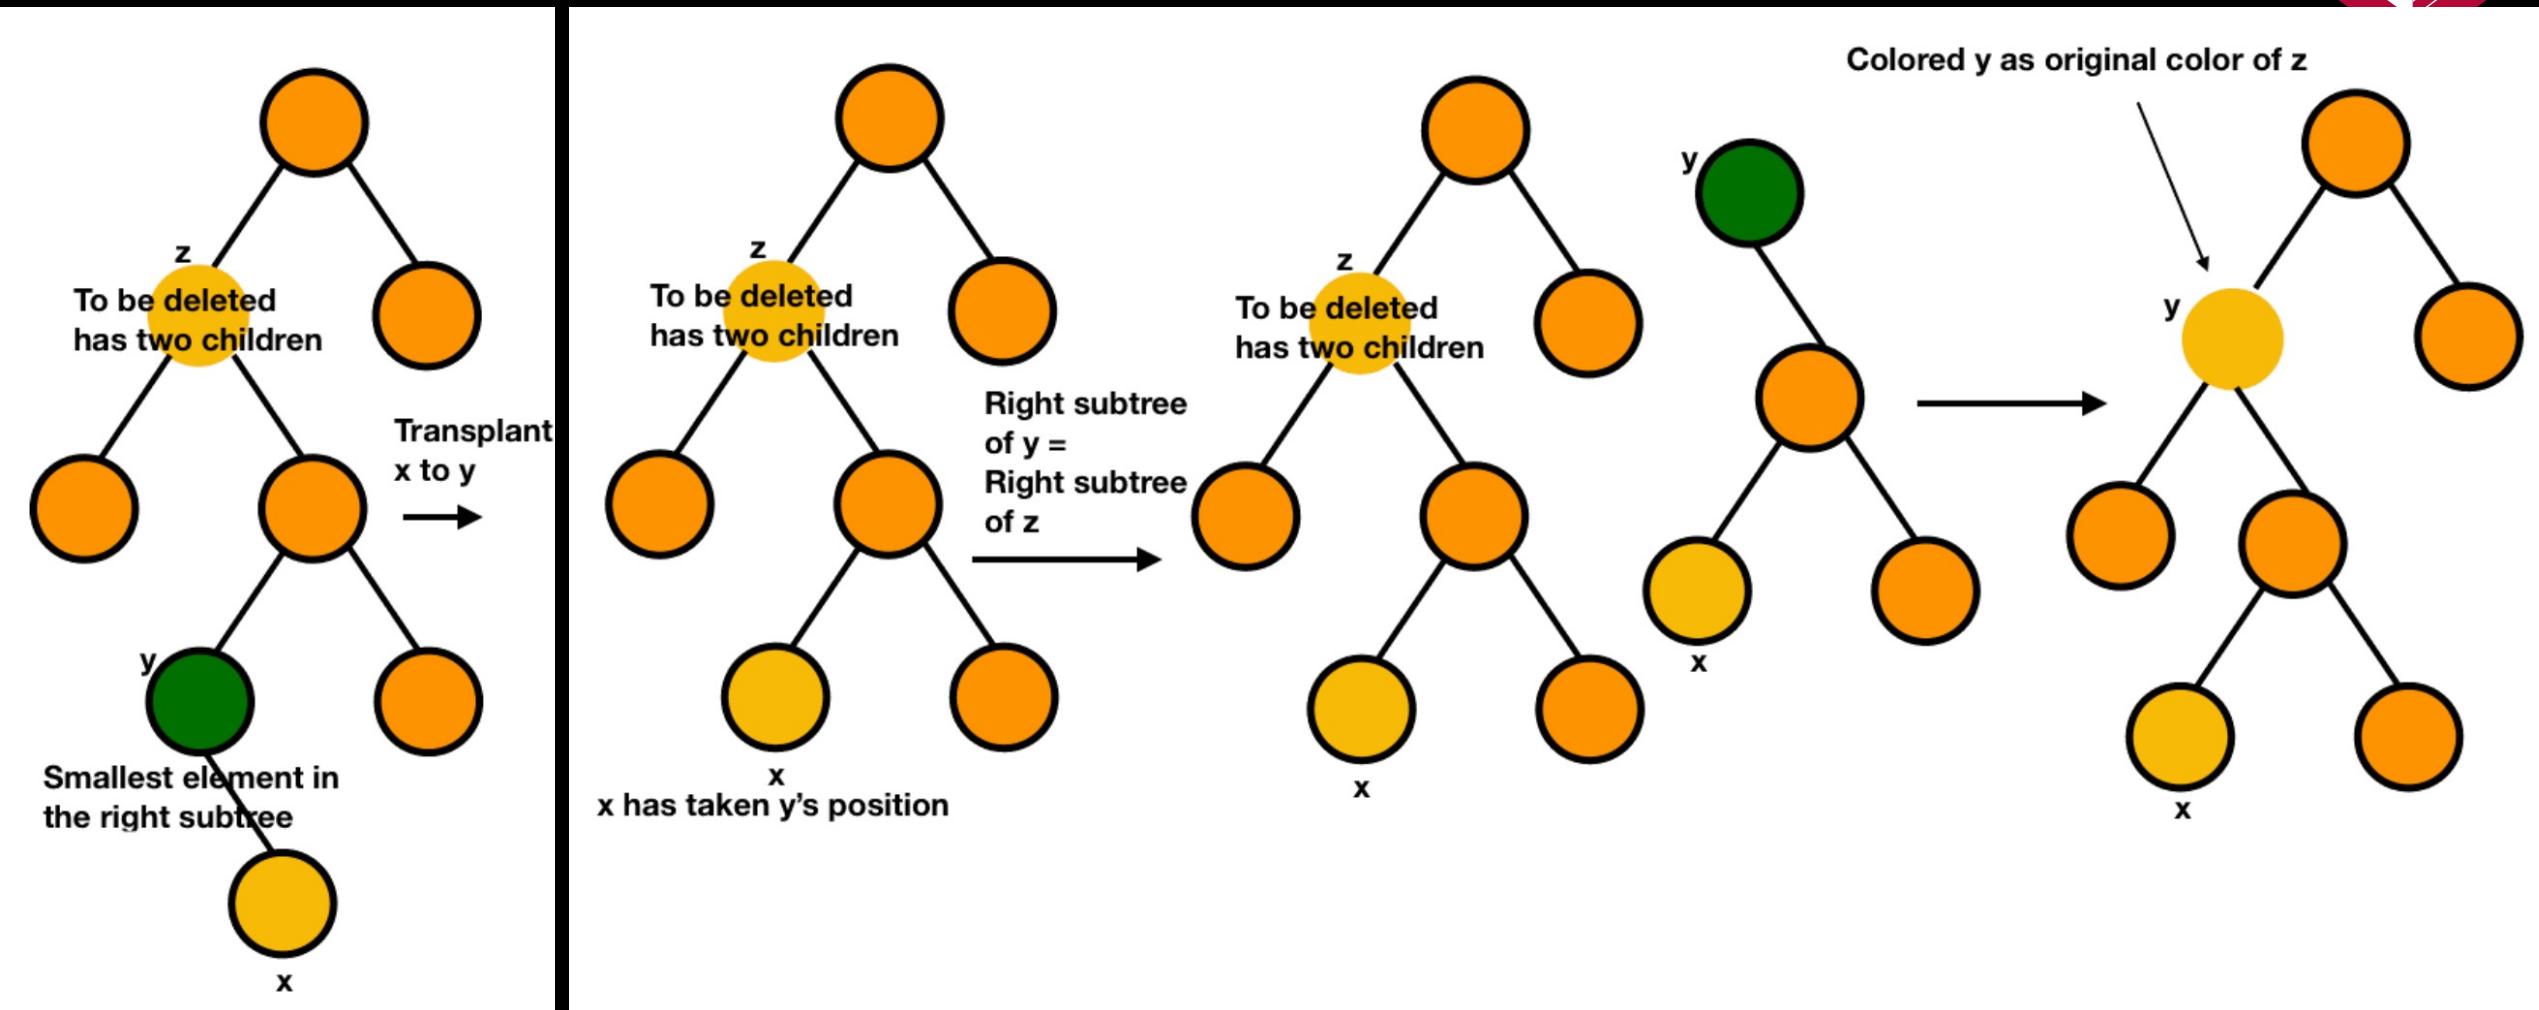
Recap – BST Deletion



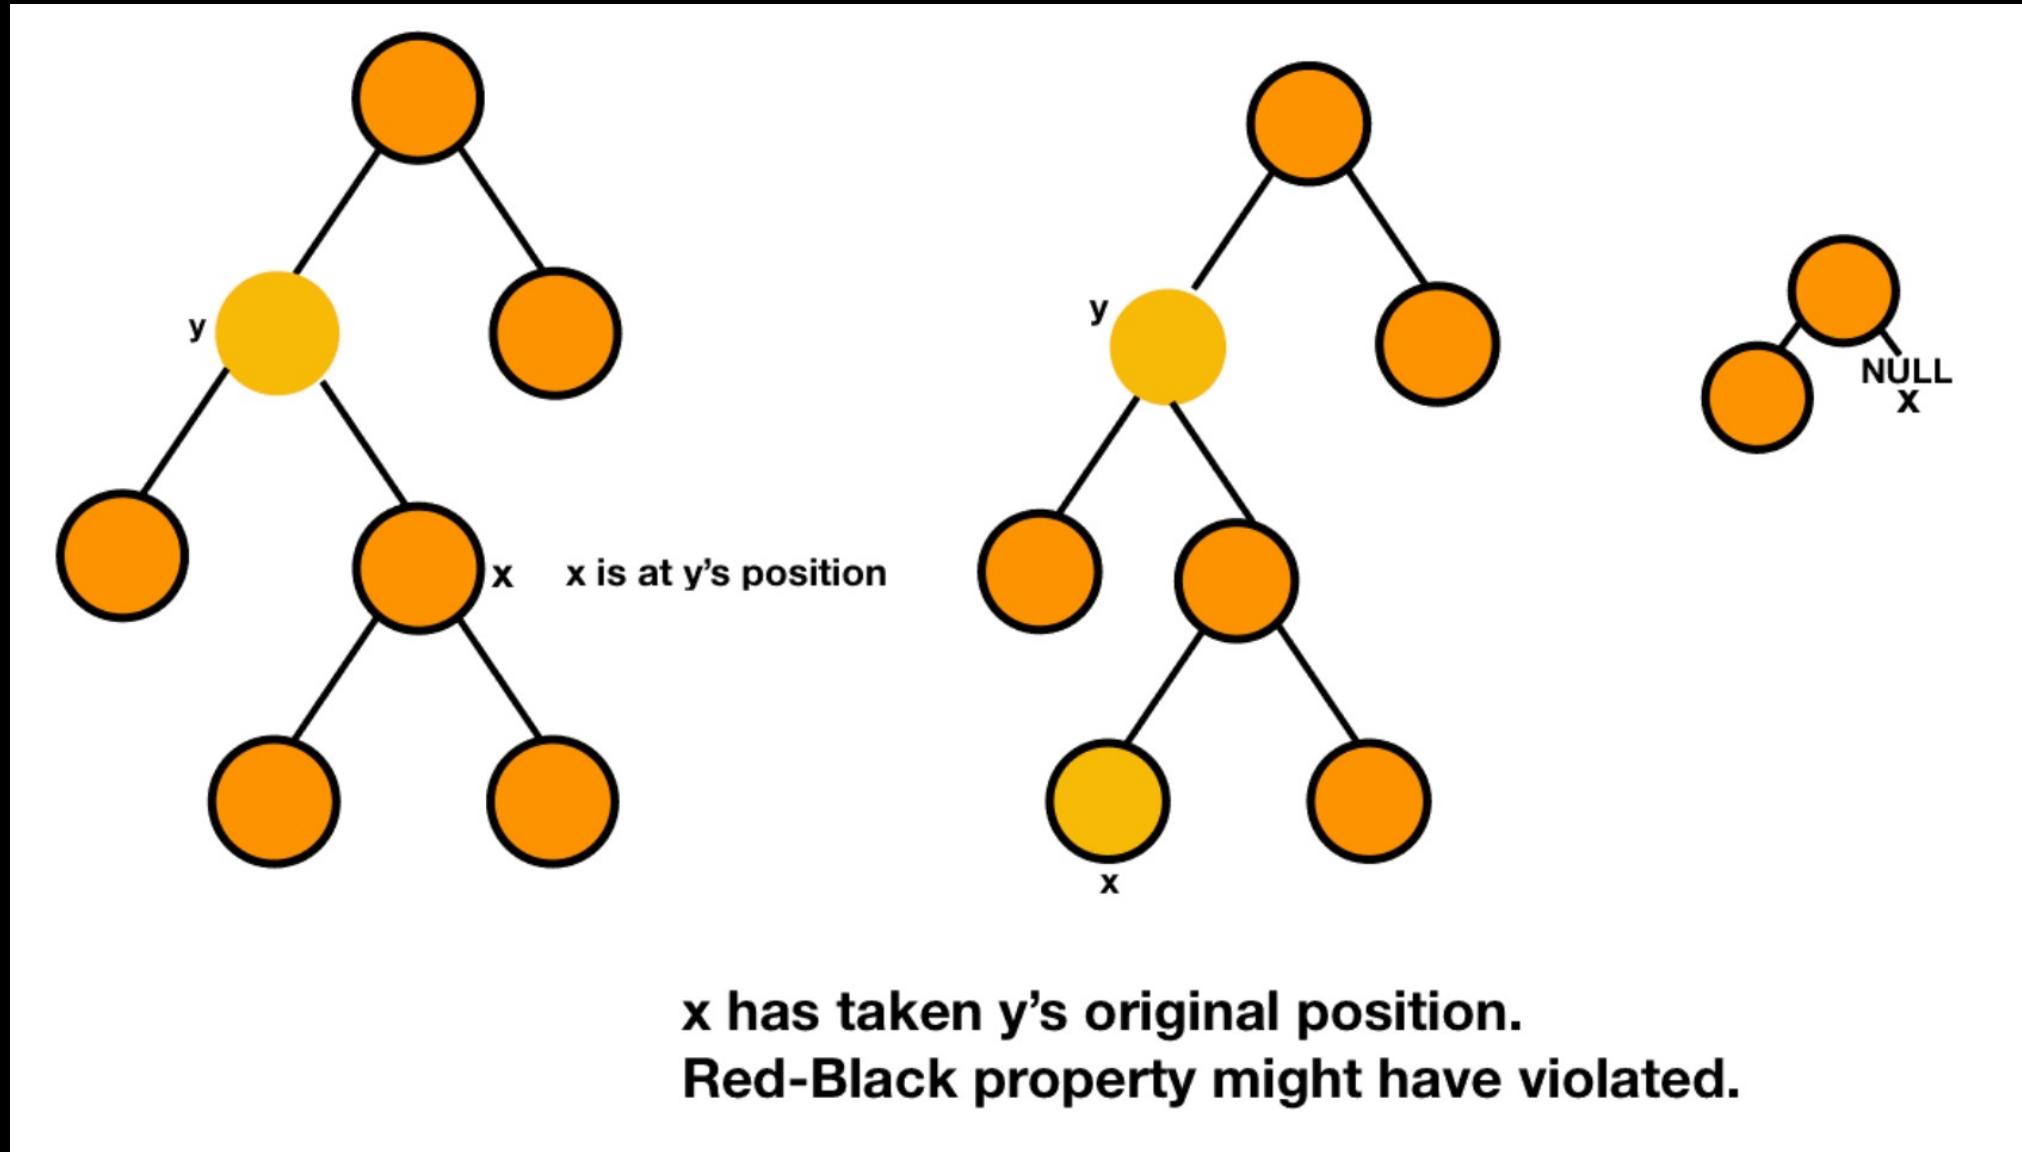
Recap – BST Deletion



Recap – BST Deletion



7.3.3. Deletions





7.3.3. Deletions

(1) RB-TRANSPLANT(T, u, v)
if $u.p = T.nil$
 $T.root \leftarrow v$
(2) elseif $u = u.p.left$
 $u.p.left \leftarrow v$
(3) else $u.p.right \leftarrow v$
 $v.p \leftarrow u.p$
(4)
(5)
(6)
(7)
(8)
(9)
(10)
(11)
(12)
(13)
(14)
(15)
(16)
(17)
(18)
(19)
(20)
(21)
(22)

RB-DELETE(T, z)
 $y = z$
 $y-original-color = y.color$
if $z.left == T.nil$
 $x = z.right$ // z has no left child
 RB-TRANSPLANT($T, z, z.right$)
elseif $z.right == T.nil$
 $x = z.left$ // just left child
 RB-TRANSPLANT($T, z, z.left$)
else $y = TREE-MINIMUM(z.right)$ //two children, y is successor of z
 $y-original-color = y.color$
 $x = y.right$
 if $y.p == z$
 $x.p = y$
 else RB-TRANSPLANT($T, y, y.right$) // y in subtree of z subtree
 $y.right = z.right$ //and not root of subtree
 $y.right.p = y$
 RB-TRANSPLANT(T, z, y) // replace z by y
 $y.left = z.left$
 $y.left.p = y$
 $y.color = z.color$
if $y-original-color == BLACK$
 RB-DELETE-FIXUP(T, x)

- Deleting a node from a red-black tree is a bit more complicated than inserting a node.
- The procedure for deleting a node from a red-black tree is based on the TREE-DELETE procedure.



7.3.3. Deletions

Difference between RB-DELETE and TREE-DELETE:

- We maintain node y as the node either removed from the tree (when z has fewer than 2 children) or moved within the tree (when z has 2 children).
- Need to save original color of y . It is tested at the end, because if it's black, then removing or moving y could cause red-black properties to be violated.
- x is the node that moves into original position of y . It is either only child of y , or $T.\text{nil}$ if y has no child.
- Set $x.p$ to point to the original position of y 's parent, even if $x=T.\text{nil}$. $x.p$ is set into one of two ways:
 - If z not original parent of y , then $x.p$ is set in last line of RB-TRANSPLANT.
 - Otherwise, y will move up to take z 's position in the tree. Assignment $x.p=y$ makes $x.p$ point to original position of y 's parent , even if x is $T.\text{nil}$.
- If original color of y was black, the changes to the tree structure might cause a violation of the red-black properties. Call RB-DELETE-FIX to resolve violations.

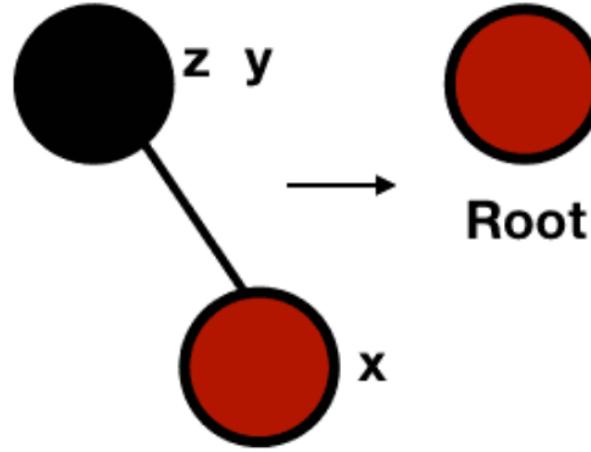


7.3.3. Deletions

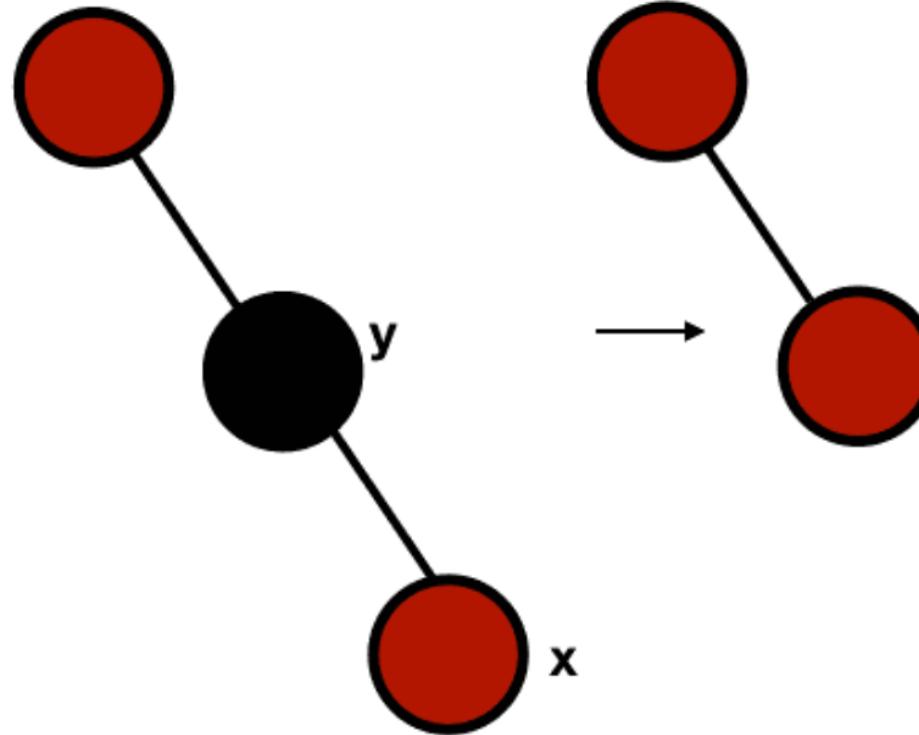
Possible violations of red-black properties:

1. No violation
2. If y is the root and x is red, then the root has become red.
3. No violation
4. Violation if x.p and x are both red.
5. Any simple path containing y now has 1 fewer black node.
 - Correct by giving x an “extra black”.
 - Add 1 to count of black nodes in paths containing x.
 - x is either doubly black (if x.color=BLACK) or red & black (if x.color=RED).
 - Property 5 satisfied, but property 1 violated.
 - The attribute x.color is still either RED or BLACK. No new values for the color attribute.
 - Extra blackness on a node is by virtue of x pointing to the node.

7.3.3. Deletions



Violation of prop. 2

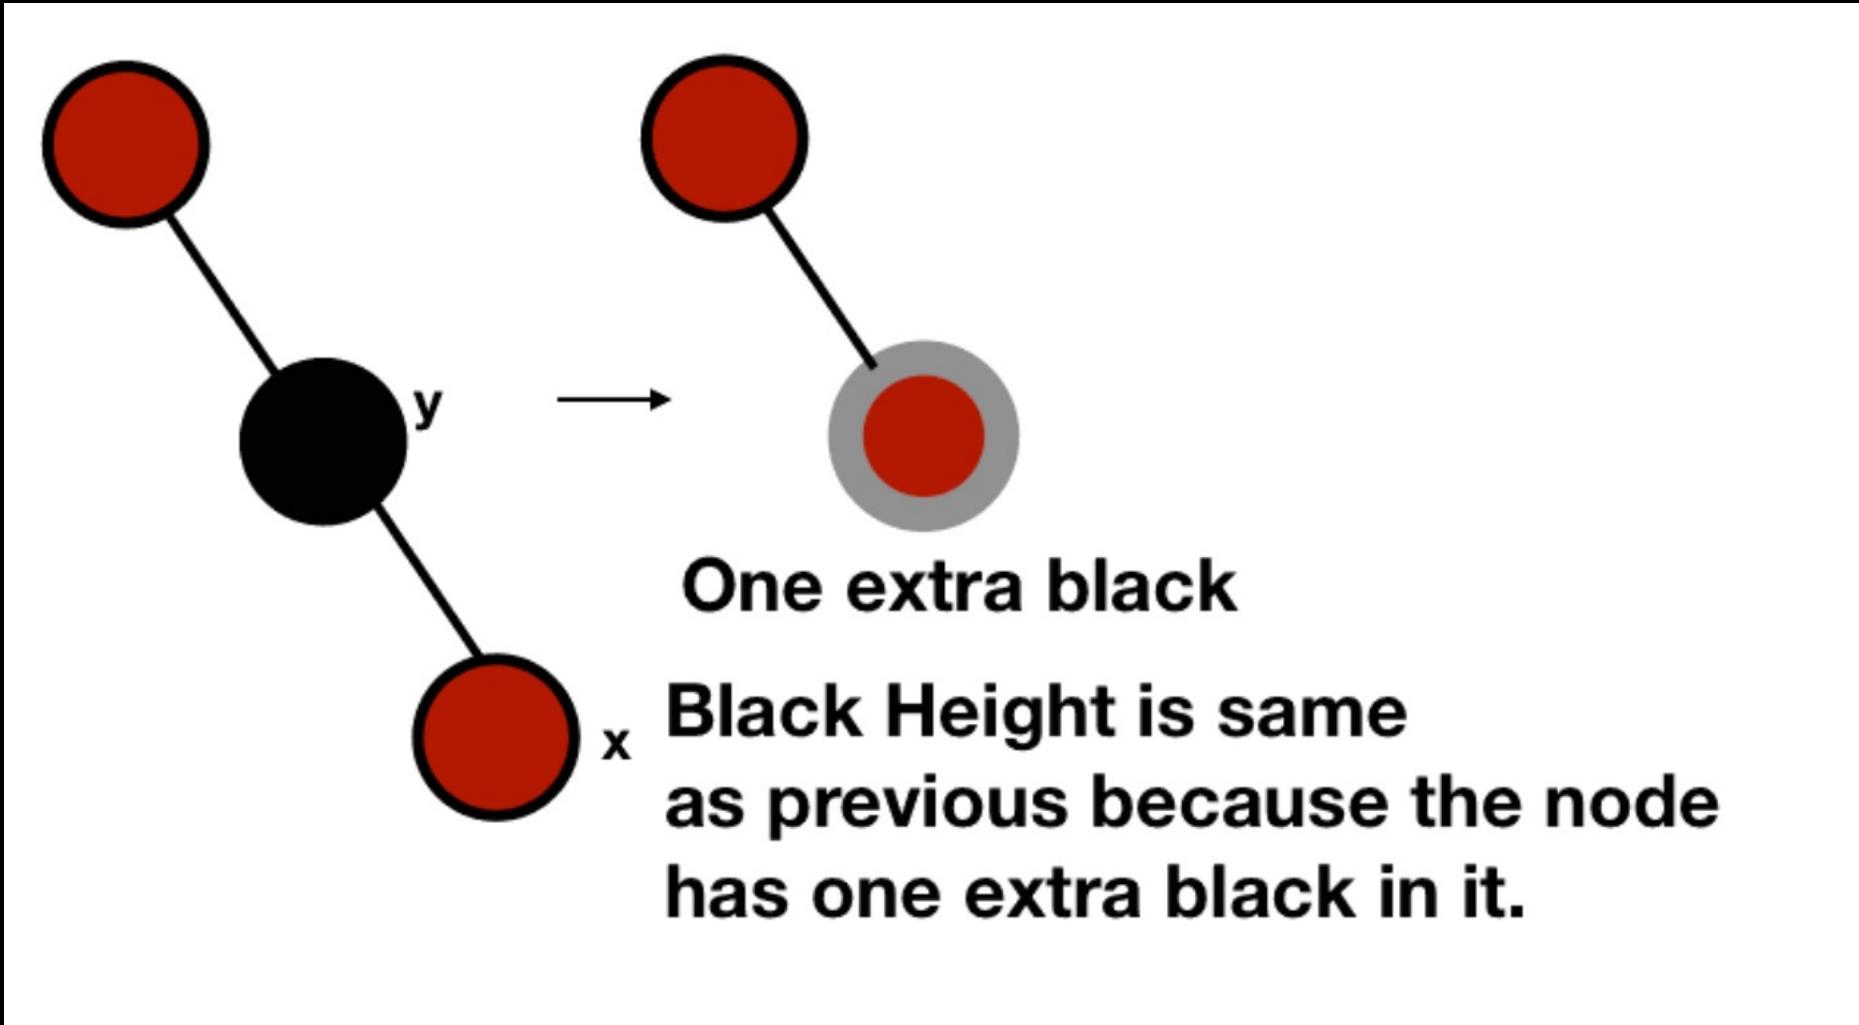


Violation of prop. 4

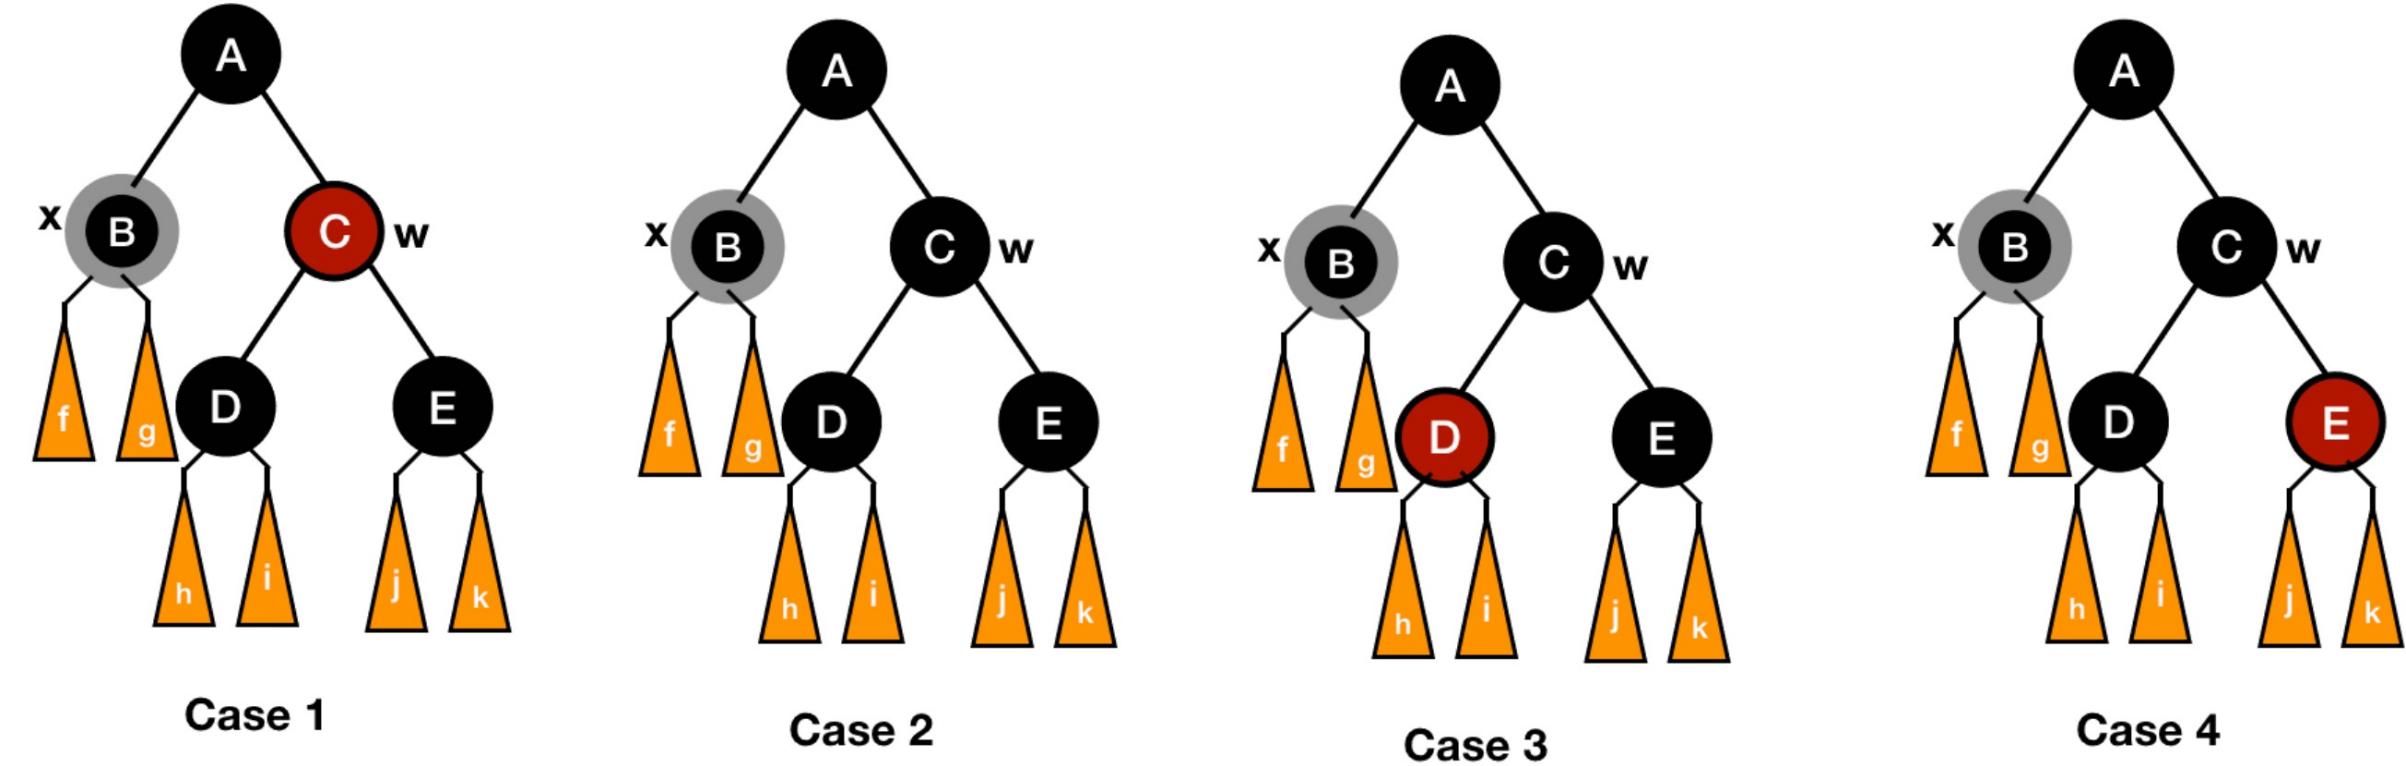
**Violation of prop. 5,
black height affected**



7.3.3. Deletions



7.3.3. Deletions



Case 1

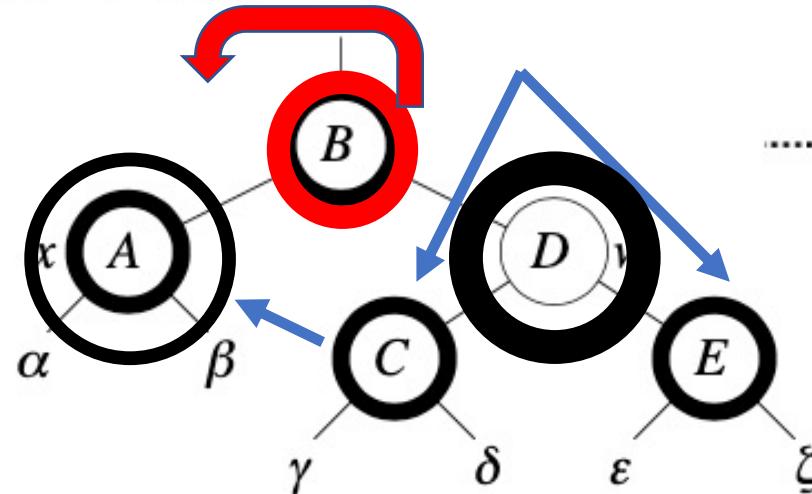
Case 2

Case 3

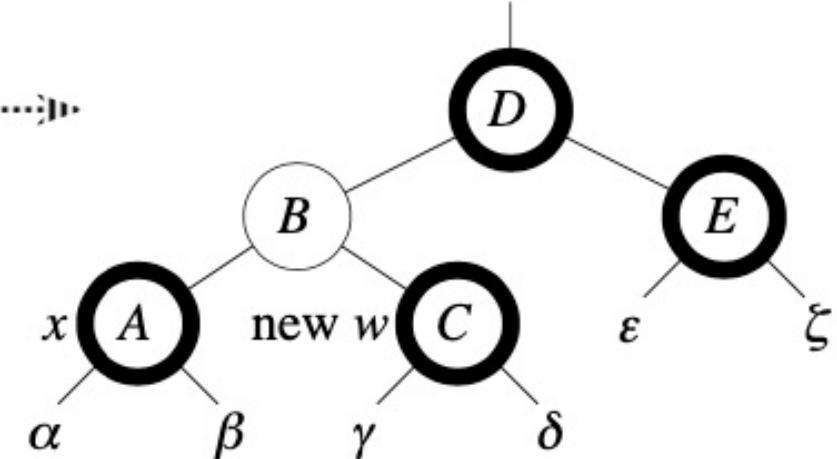
Case 4

7.3.3. Deletions

Case 1: w is red



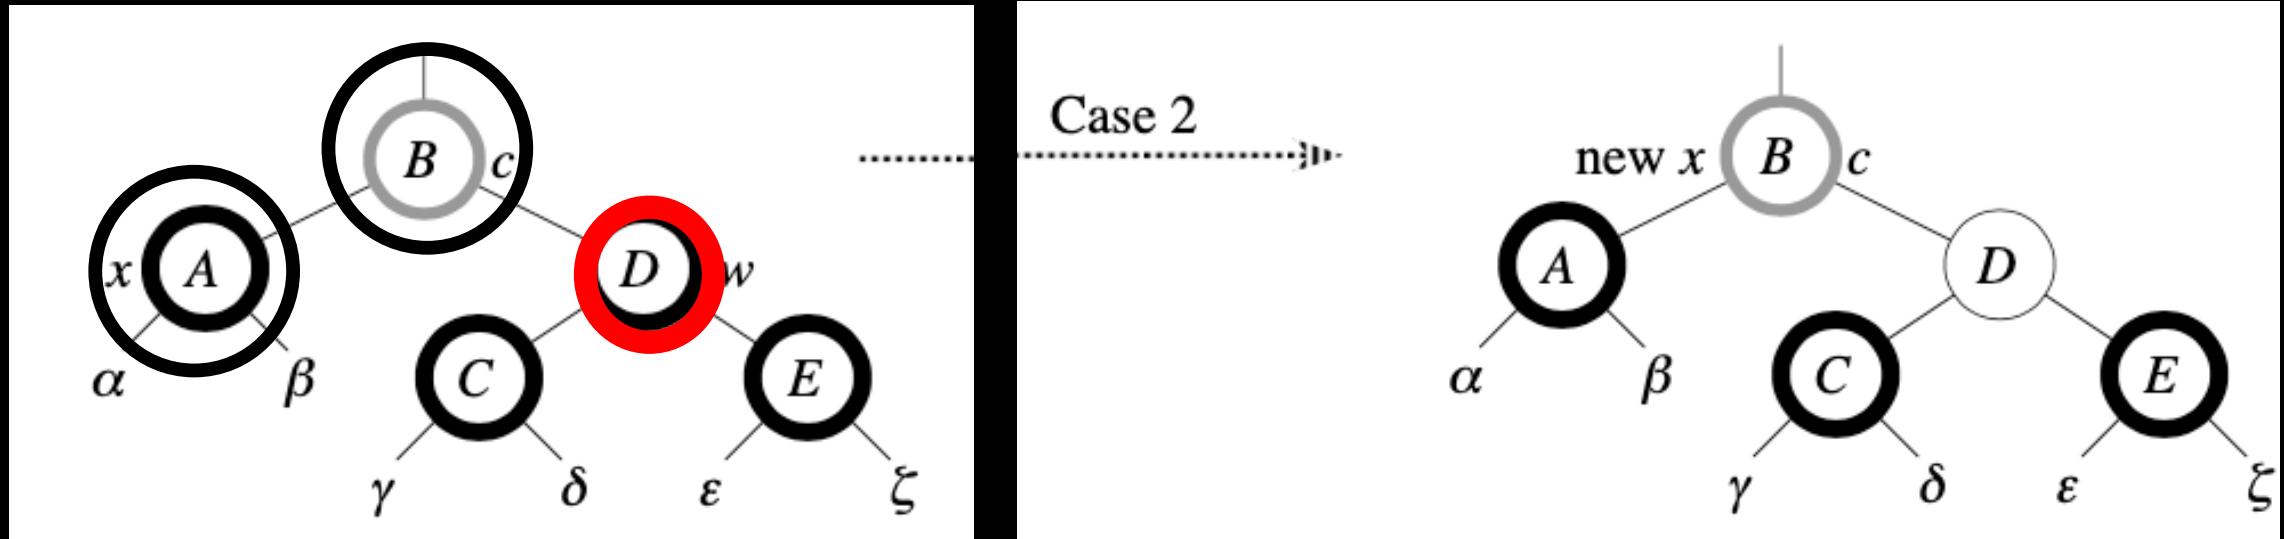
Case 1



- **w must have black children.**
- Make w black and $x.p$ red.
- Then left rotate on $x.p$.
- **New sibling of x was a child of w before rotation \Rightarrow must be black.**
- Go immediately to case 2, 3, or 4.

7.3.3. Deletions

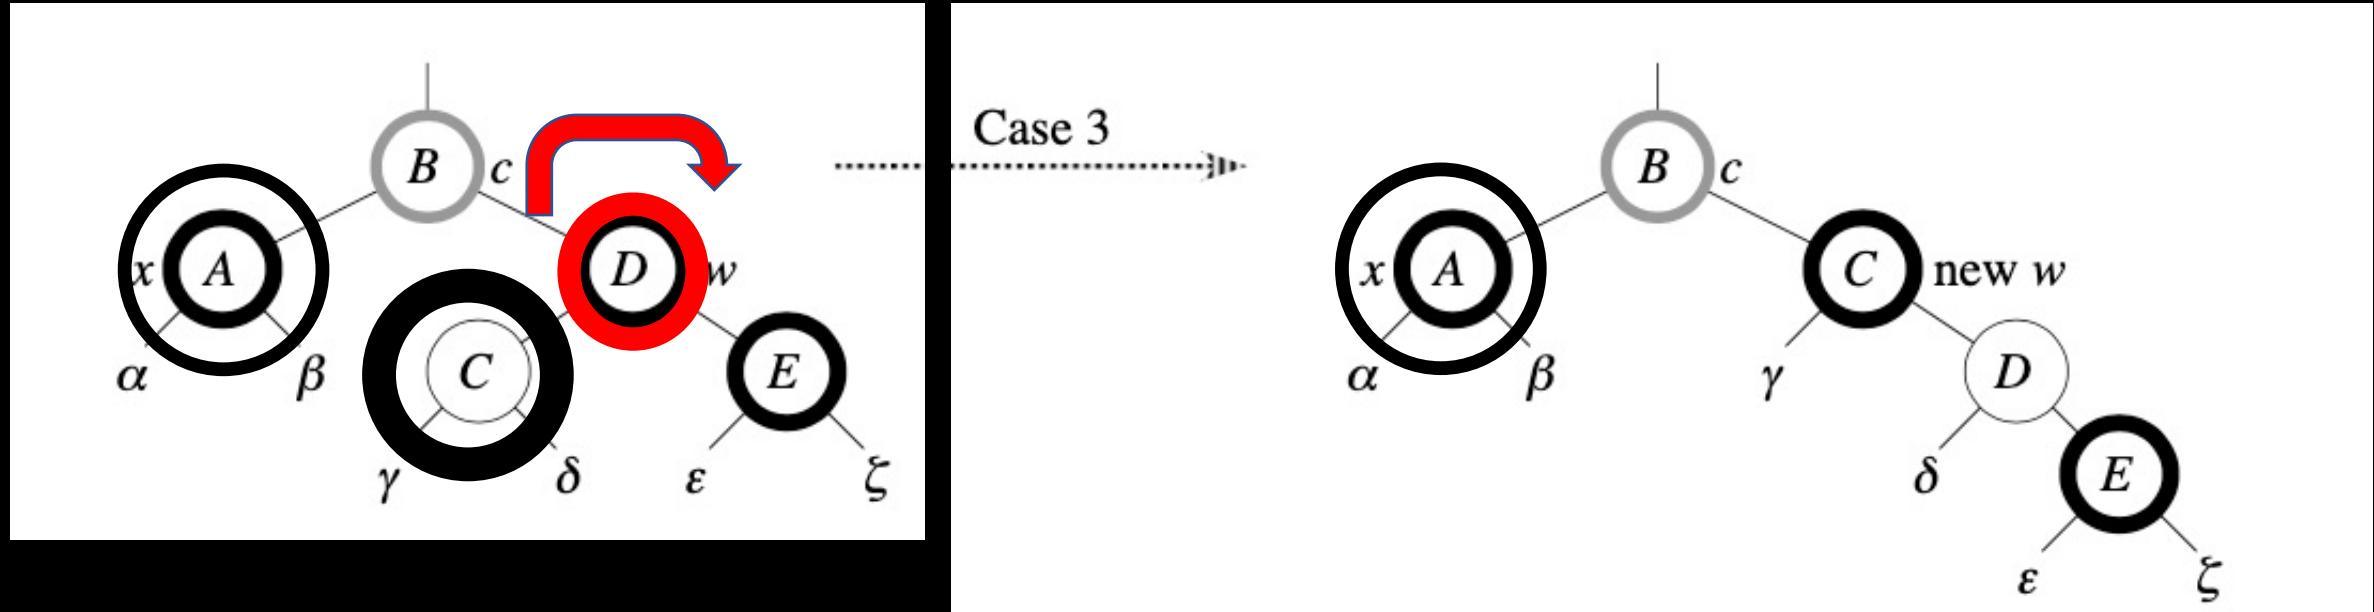
Case 2: w is black and both of w 's children are black



- Take 1 black off x (\Rightarrow singly black) and off w (\Rightarrow red).
- Move that black to $x.p$.
- Do the next iteration with $x.p$ as the new x .
- If entered this case from case 1, then $x.p$ was red \Rightarrow new x is red & black \Rightarrow color attribute of new x is RED \Rightarrow loop terminates.
 - Then new x is made black in the last line.

7.3.3. Deletions

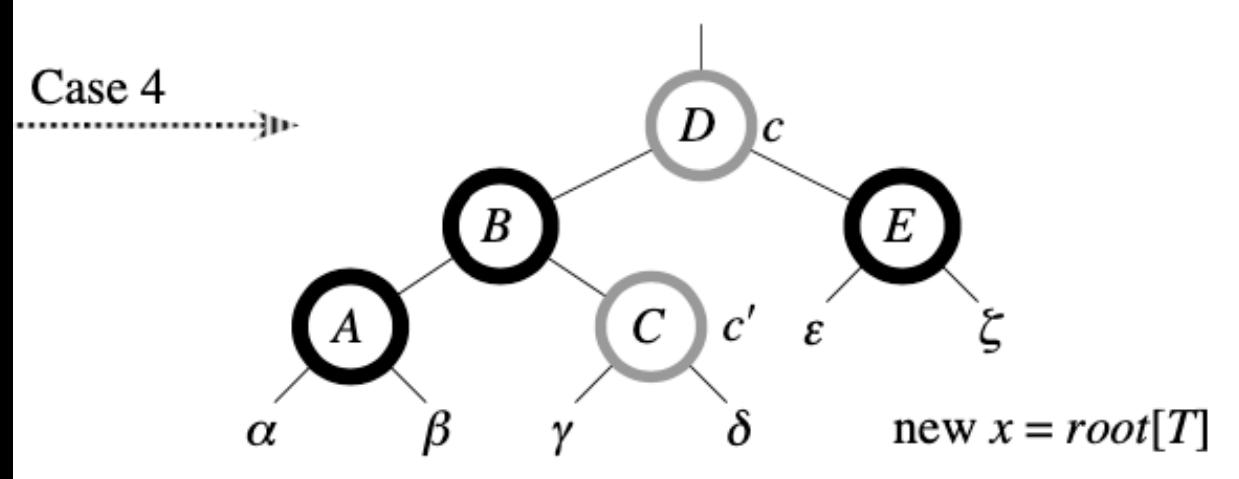
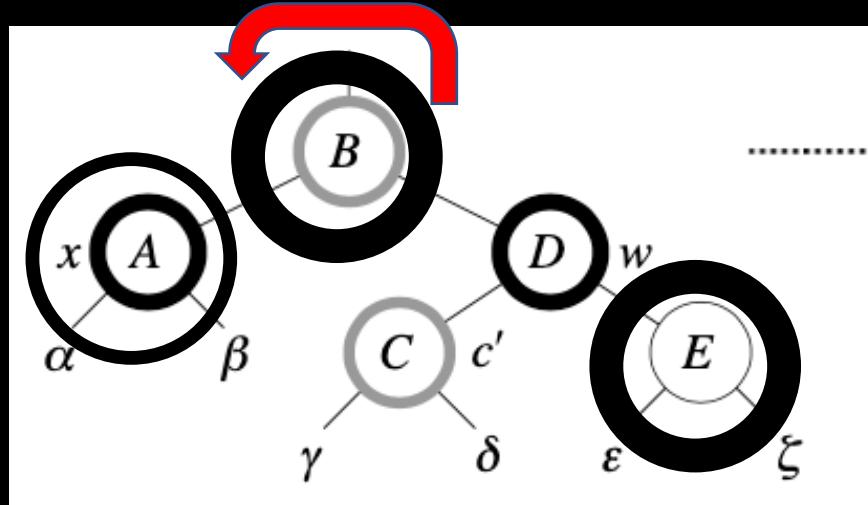
Case 3: w is black, w 's left child is red, and w 's right child is black



- Make w red and w 's left child black.
- Then right rotate on w .
- New sibling w of x is black with a red right child \Rightarrow case 4.

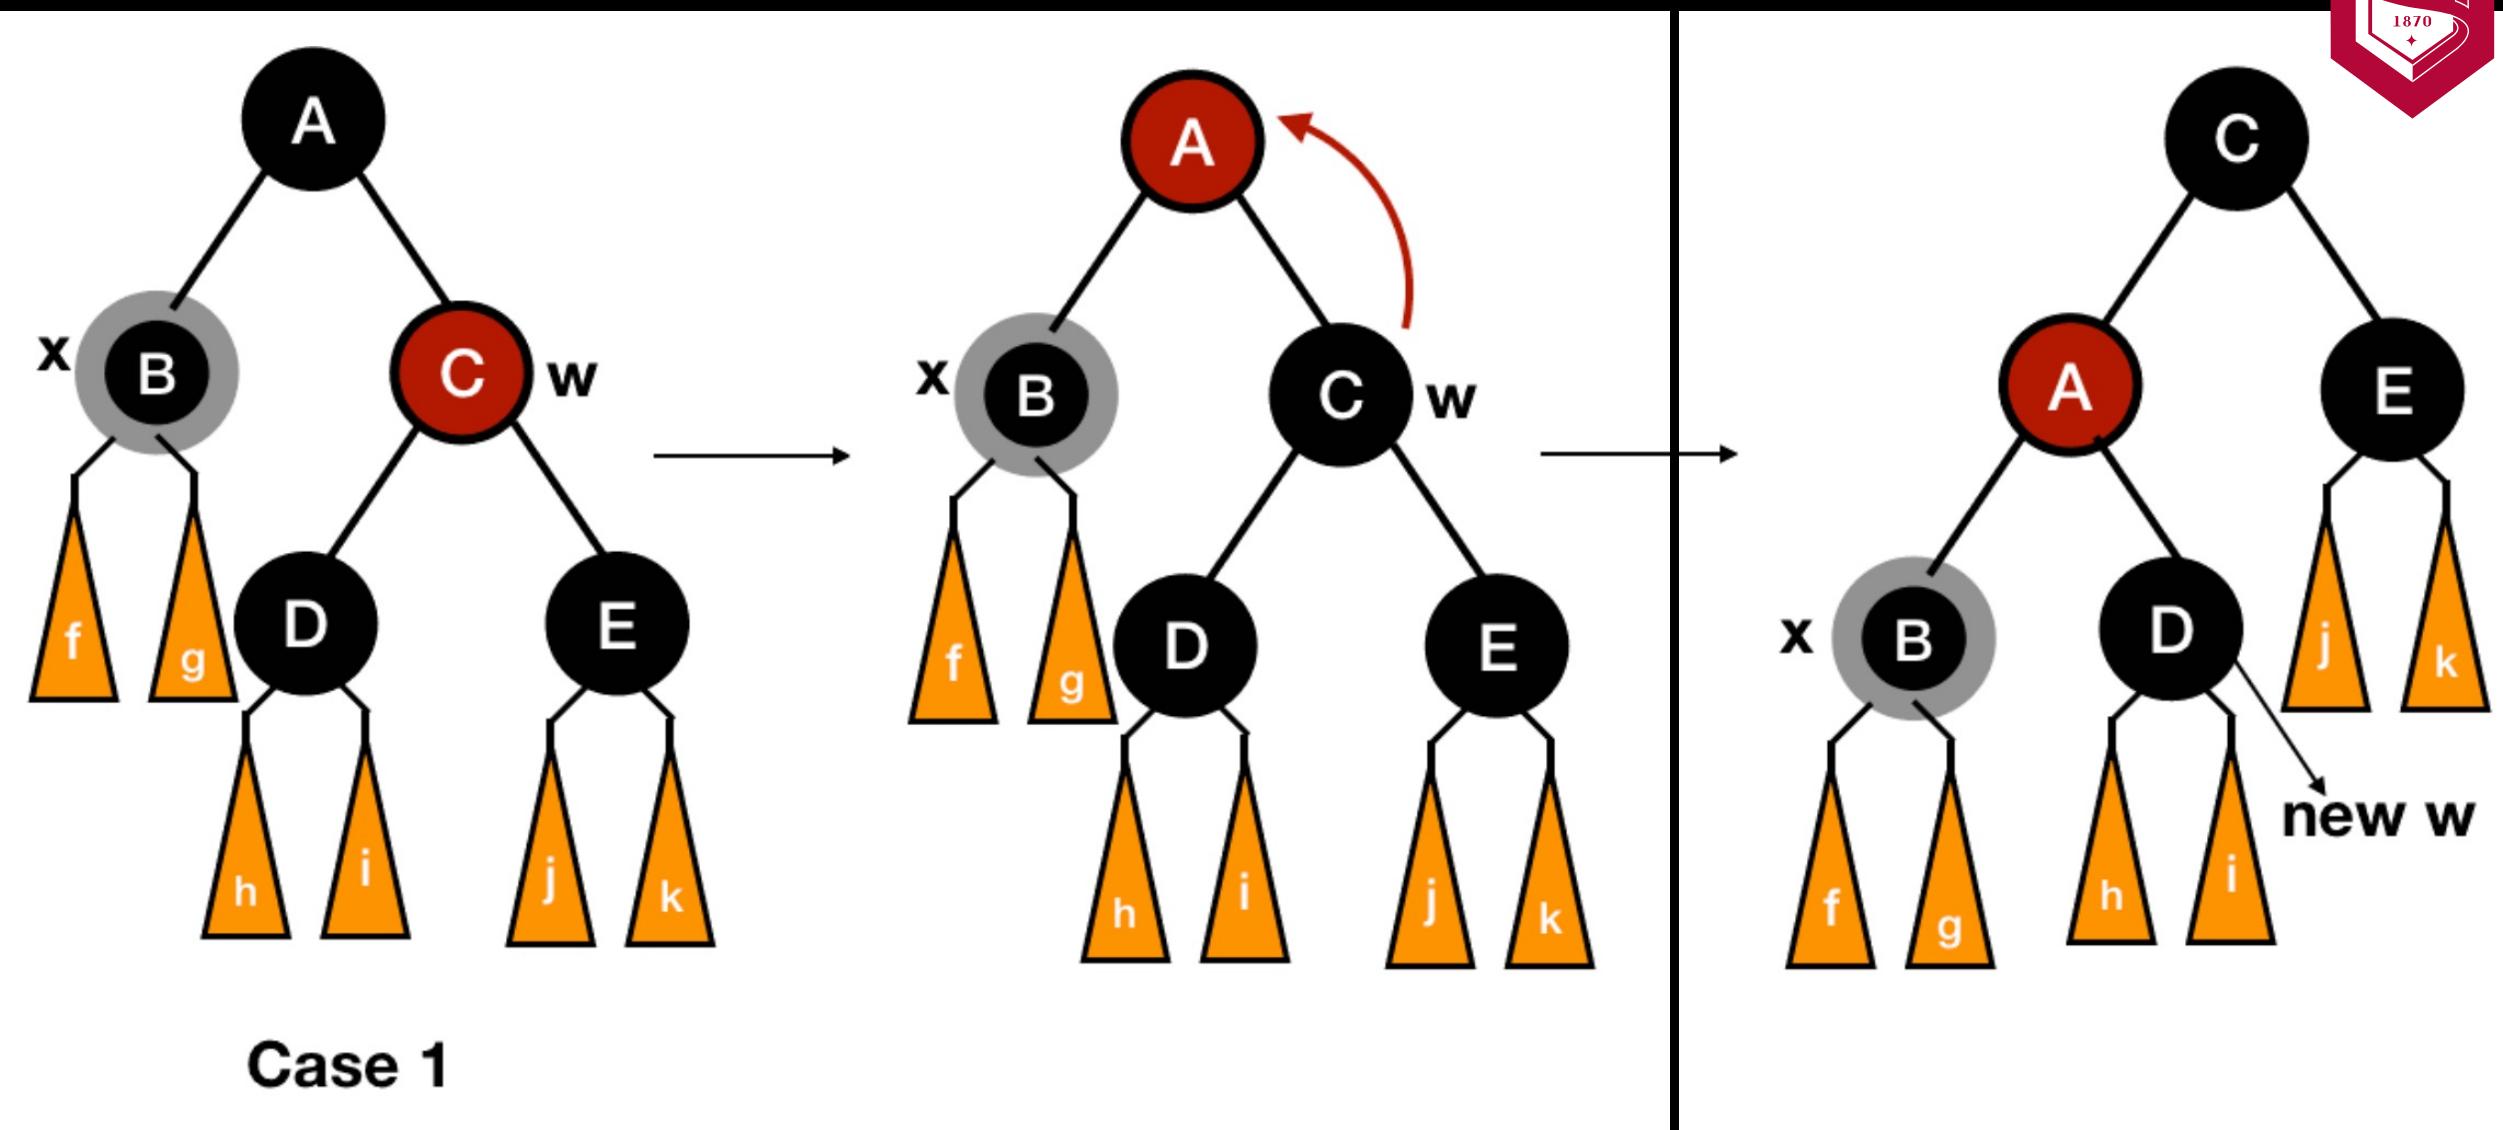
7.3.3. Deletions

Case 4: w is black, w 's left child is black, and w 's right child is red

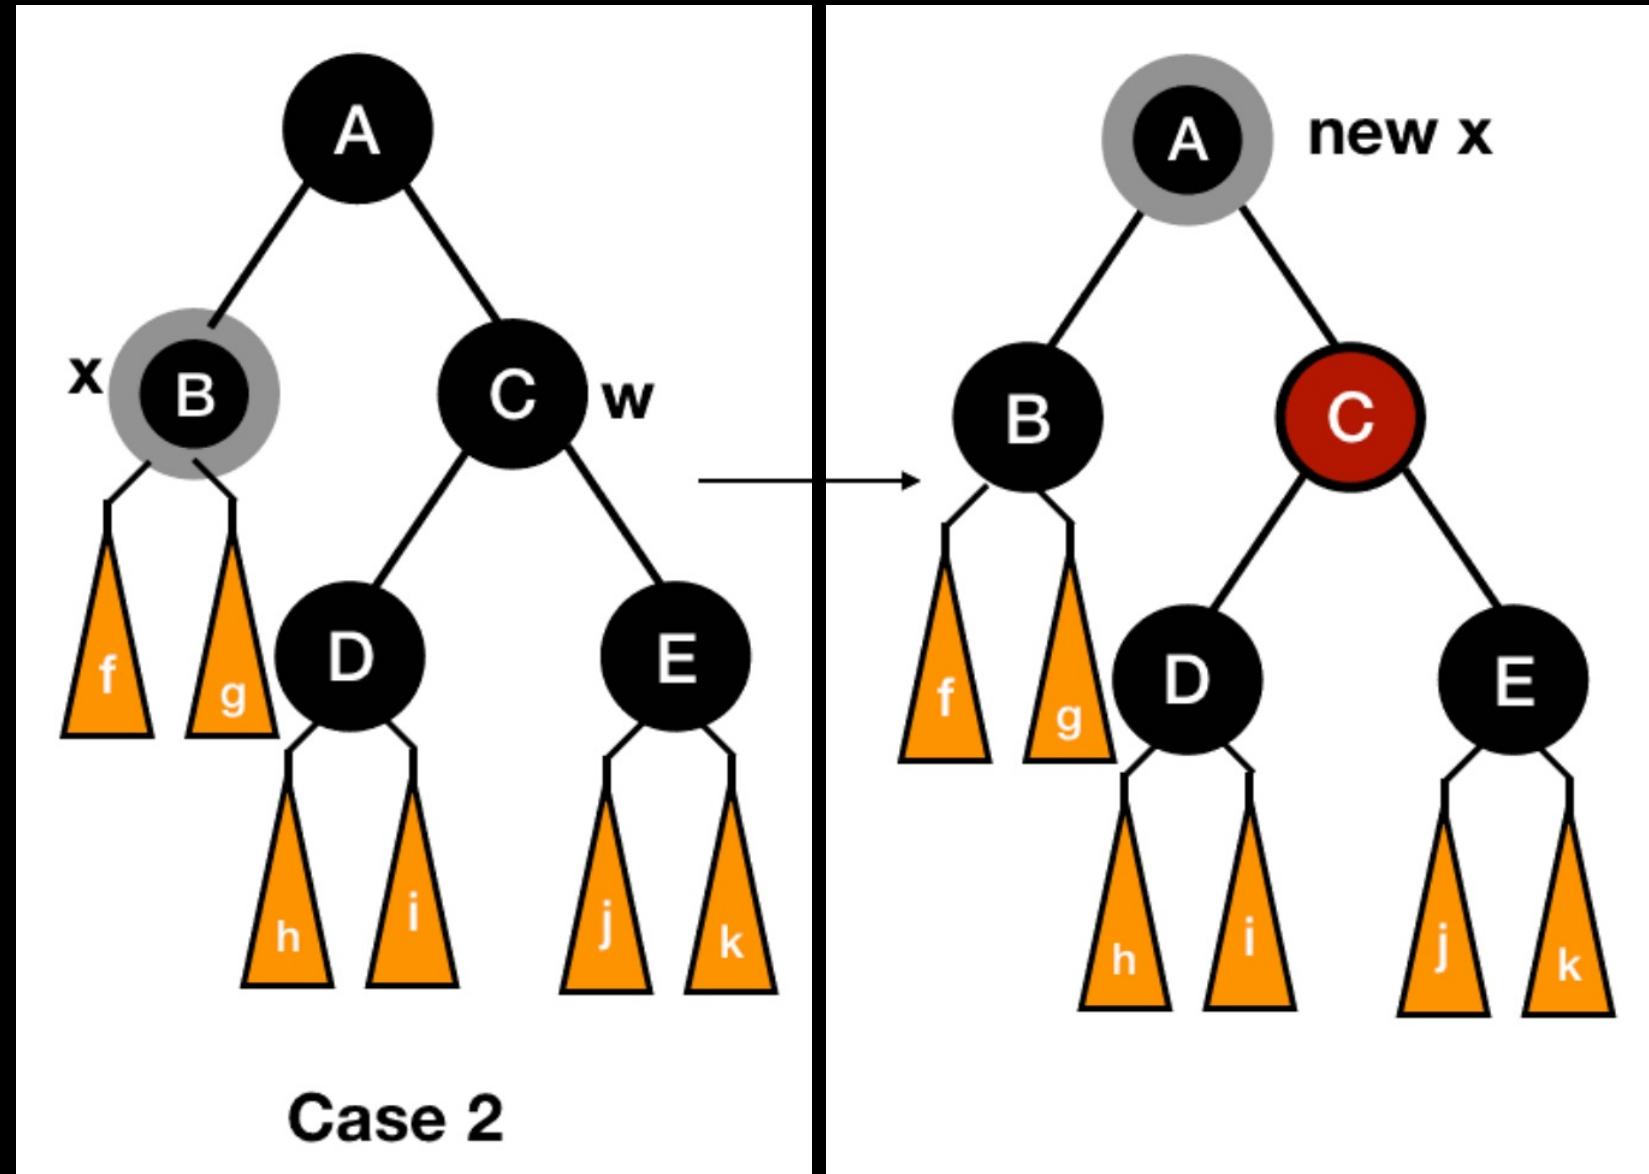


- Make w be $x.p$'s color (c).
- Make $x.p$ black and w 's right child black.
- Then left rotate on $x.p$.
- Remove the extra black on x ($\Rightarrow x$ is now singly black) without violating any red-black properties.
- All done. Setting x to root causes the loop to terminate.

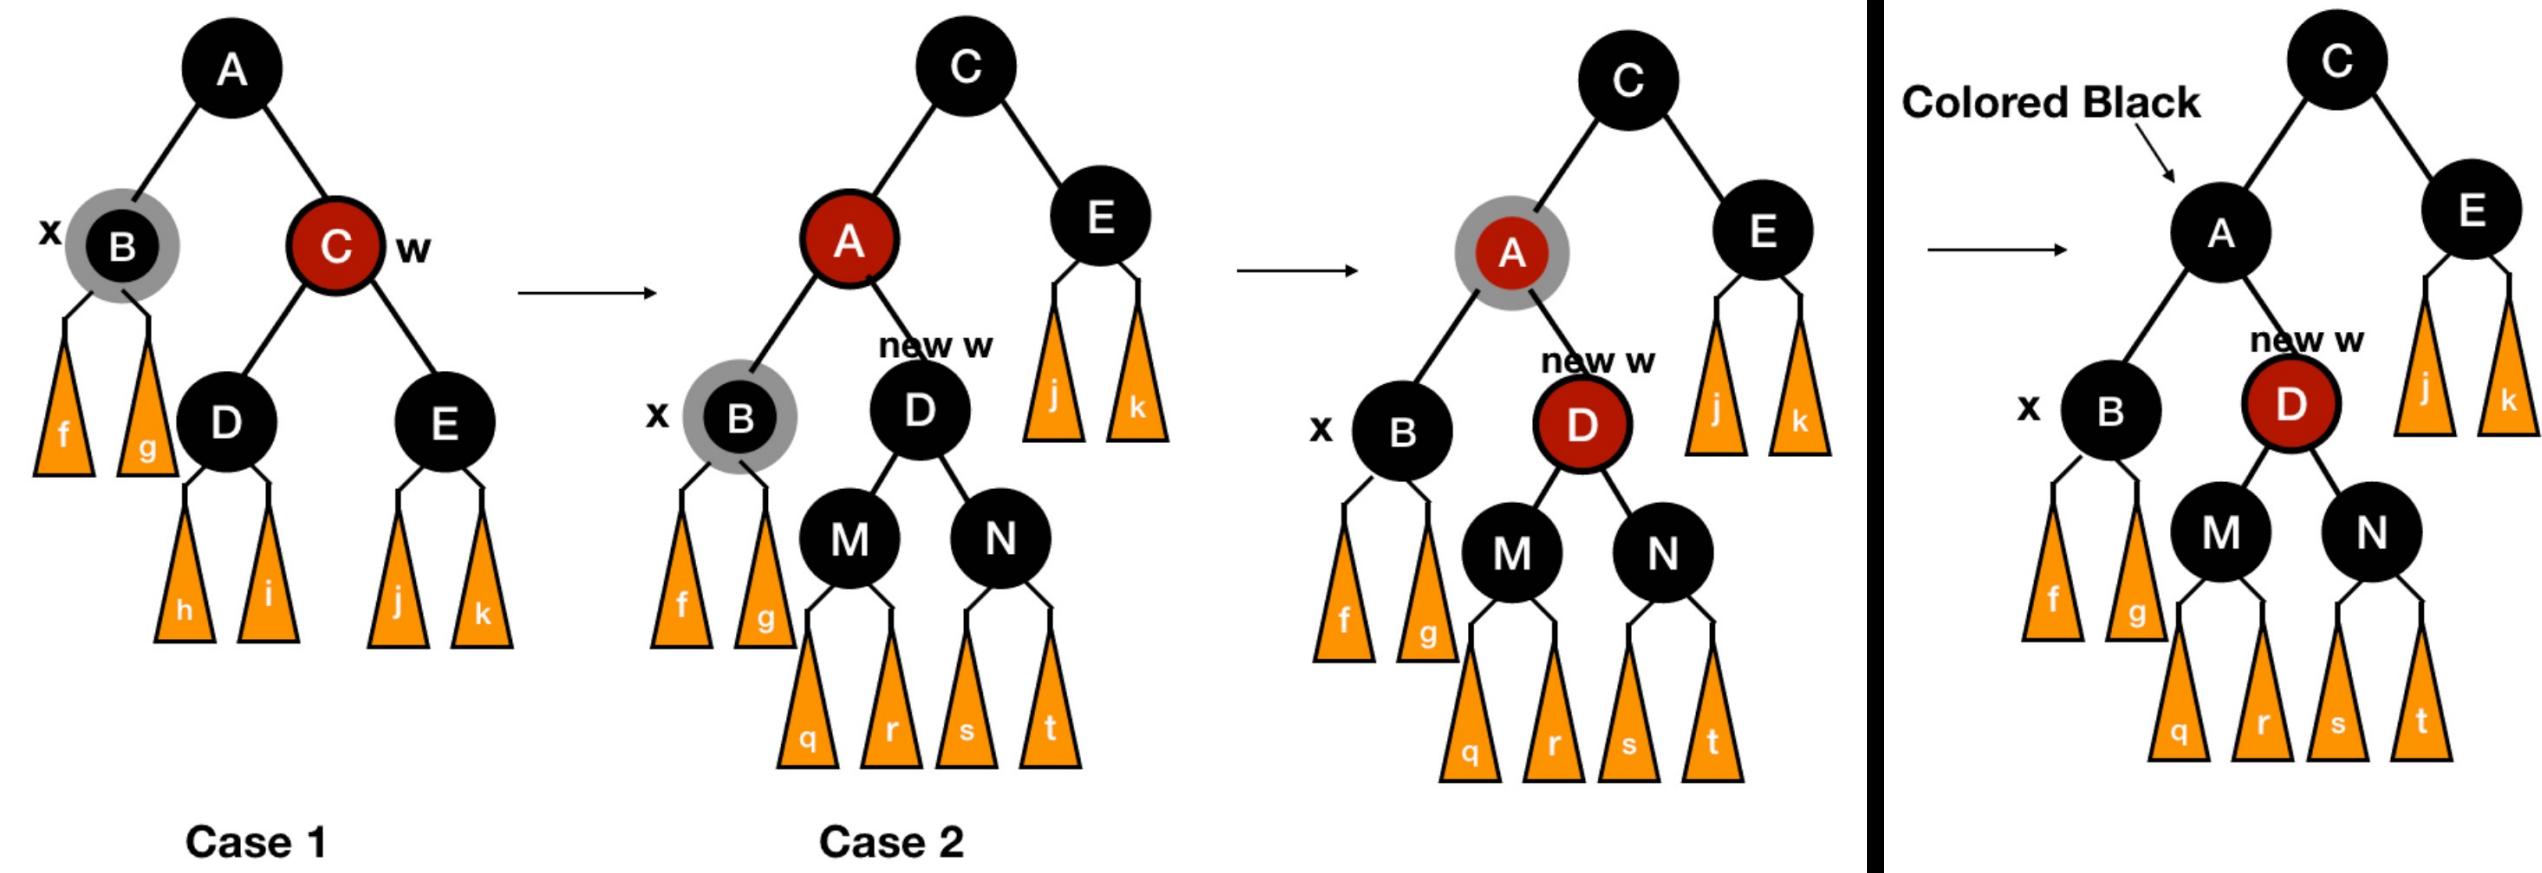
7.3.3. Deletions



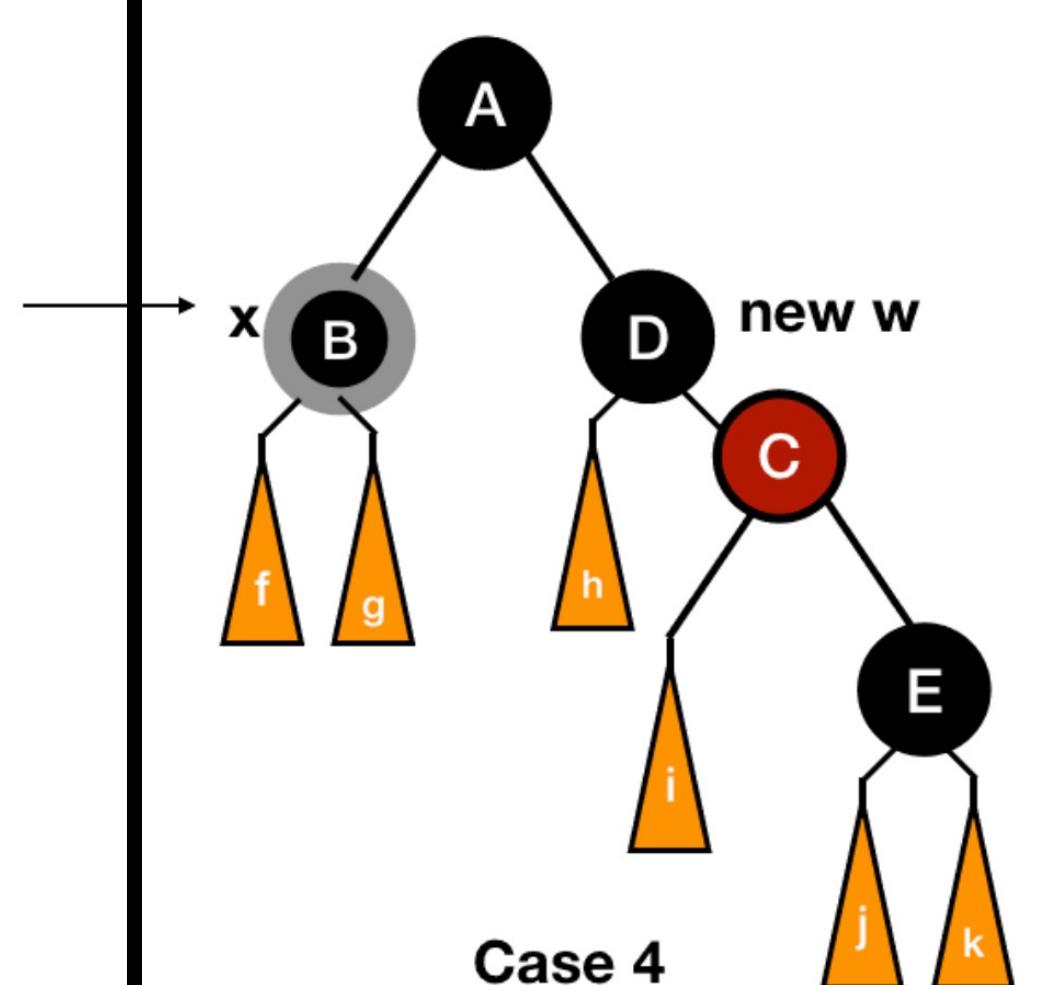
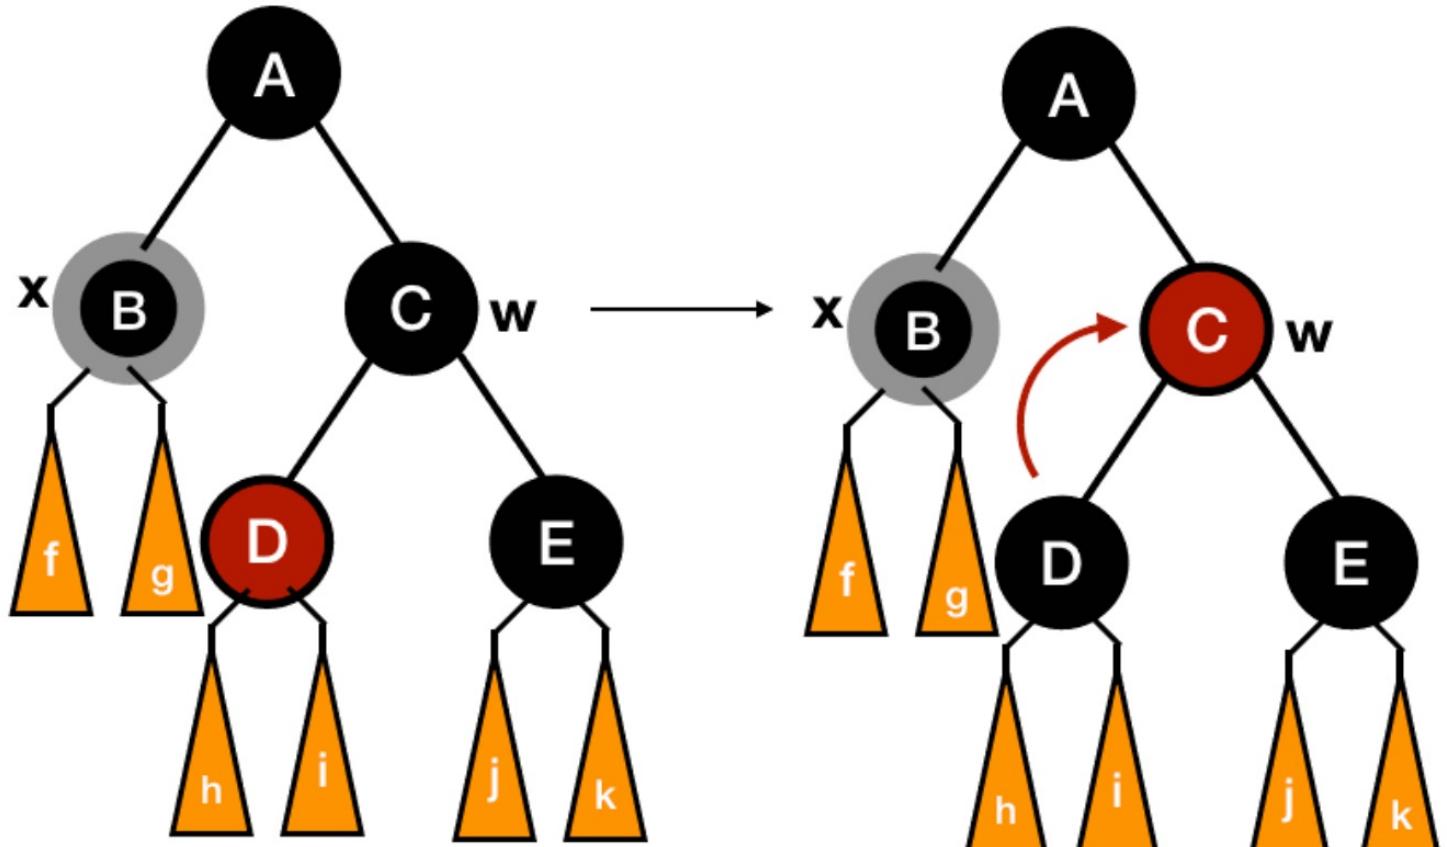
7.3.3. Deletions



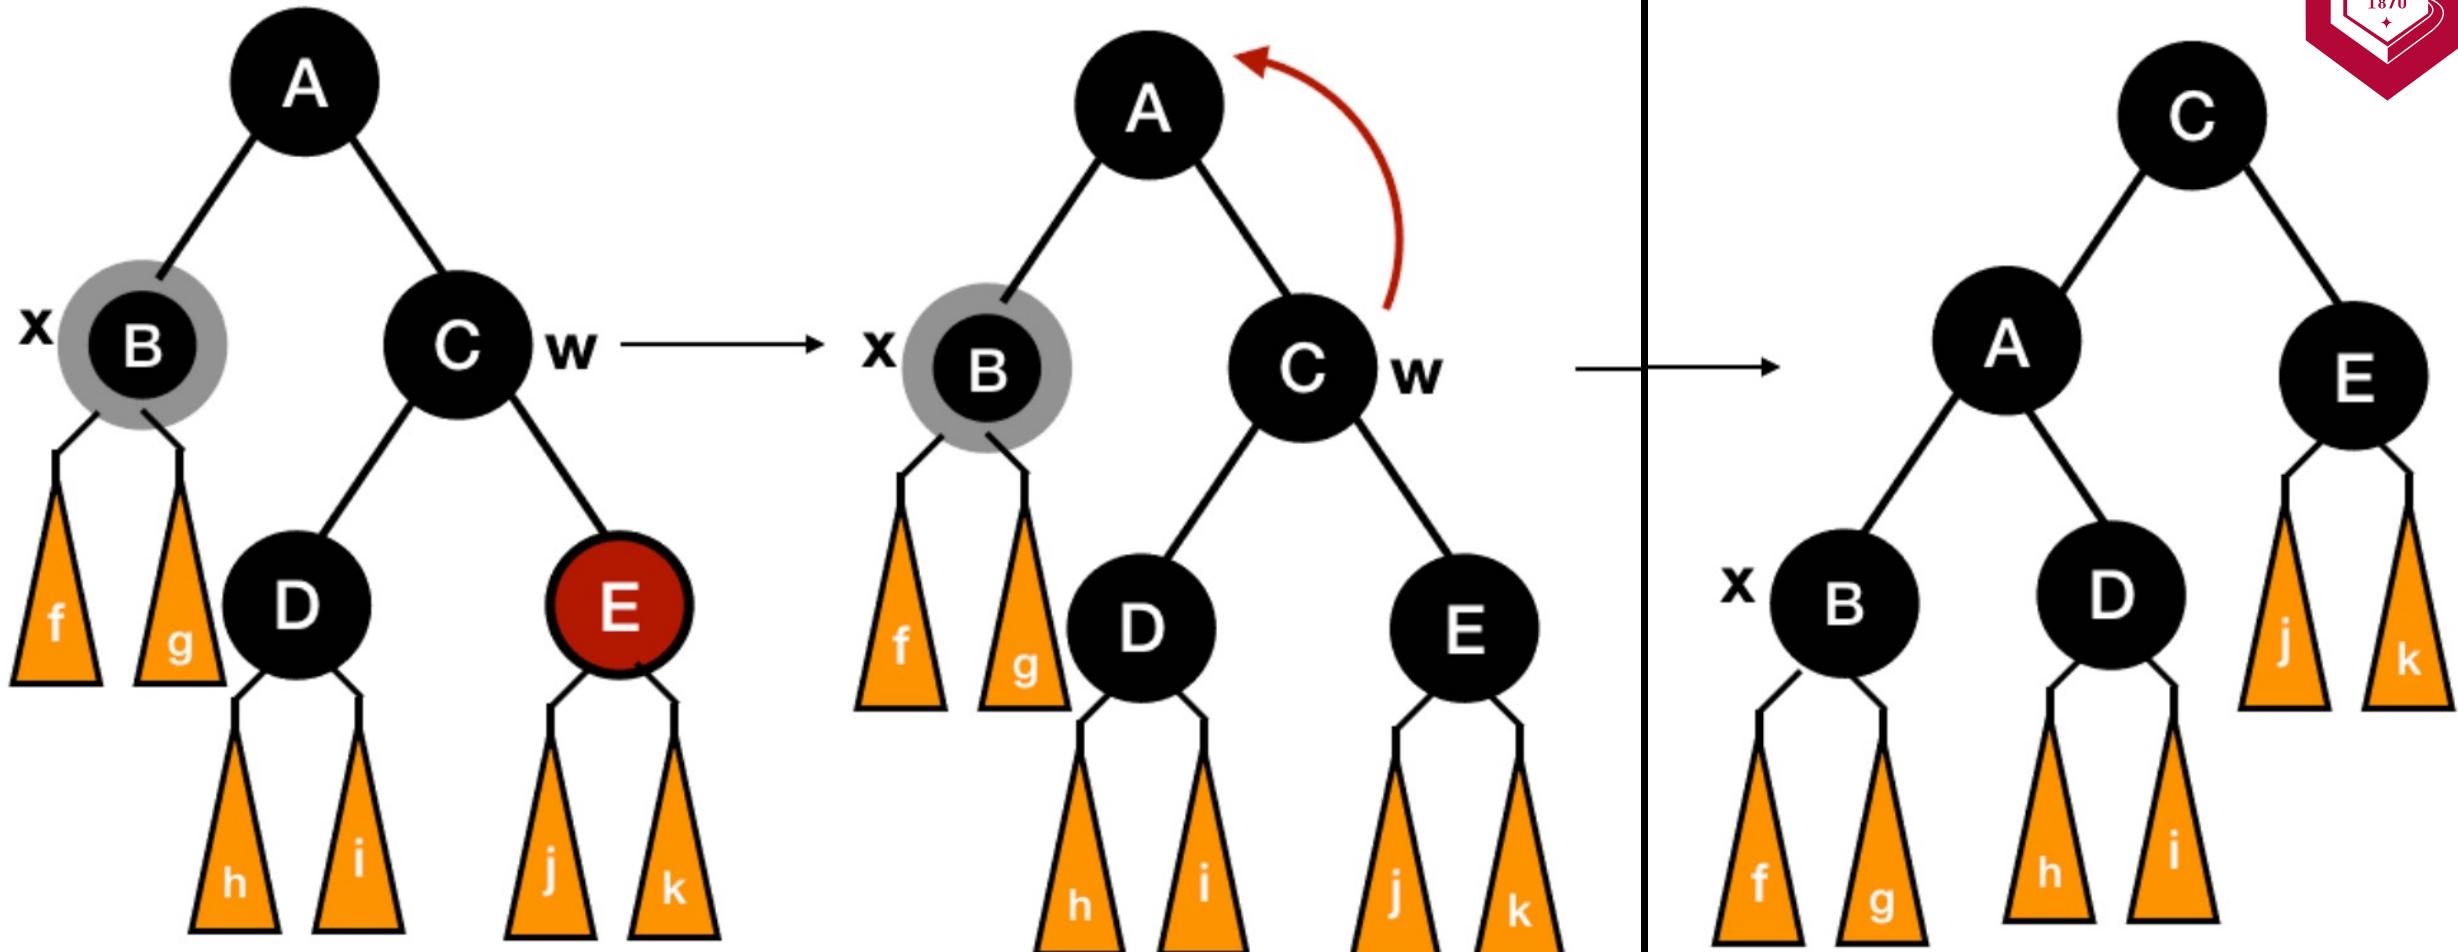
7.3.3. Deletions



7.3.3. Deletions



7.3.3. Deletions



Case 4



7.3.3. Deletions

Idea: Move the extra black up the tree until

- x points to a red & black node \Rightarrow turn it into a black node,
- x points to the root \Rightarrow just remove the extra black, or
- we can do certain rotations and re-colorings and finish.

Within the **while** loop:

- x always points to a non-root doubly black node.
- w is x 's sibling.
- w cannot be $T.nil$, since that would violate property 5 at $x.p$.

There are 8 cases, 4 of which are symmetric to the other 4. As with insertion, the cases are not mutually exclusive. We'll look at cases in which x is a left child.

```
(1) RB-DELETE-FIXUP(T,x)
(2) while ( $x \neq T.root$  and  $x.color == BLACK$ ) do
(3)   if ( $x == x.p.left$ ) then
(4)     w =  $x.p.right$ 
(5)     if ( $w.color == RED$ ) then
(6)       w.color == BLACK                                //case 1
(7)       x.p.color = RED
(8)       LEFT-ROTATE(T, x.p)
(9)     w =  $x.p.right$ 
(10)    if ( $w.left.color == BLACK$  and  $w.right.color == BLACK$ ) then
(11)      w.color = RED                                 //case 2
(12)      x = x.p
(13)    else if ( $w.right.color == BLACK$ ) then
(14)      w.left.color = BLACK                         //case 3
(15)      w.color = RED
(16)      RIGHT-ROTATE(T,w)
(17)      w =  $x.p.right$ 
(18)    w.color = x.p.color                          //case 4
(19)    x.p.color = BLACK
(20)    w.right.color = BLACK
(21)    LEFT-ROTATE(T, x.p)
(22)    x = T.root
(23)  else (same as then clause with "right" and "left" exchanged)
      x.color = BLACK
```



7.3.3. Deletions

Analysis

$O(\lg n)$ time to get through RB-DELETE up to the call of RB-DELETE-FIXUP.

Within RB-DELETE-FIXUP:

1. Case 2 is the only case in which more iterations occur.
 - x moves up 1 level.
 - Hence, $O(\lg n)$ iterations.
2. Each of cases 1, 3, and 4 has 1 rotation $\Rightarrow \leq 3$ rotations in all.
3. Hence, $O(\lg n)$ time.